



PostgreSQL

Advanced Postgres

Surendra Panpaliya

Agenda

Introduction to PostgreSQL

History and development of PostgreSQL

Key features and benefits of PostgreSQL

PostgreSQL vs. Oracle database systems

Agenda



PostgreSQL vs Oracle: Key architectural differences



PostgreSQL installation & tools (pgAdmin, DBeaver)



Overview of pg_catalog, schemas, roles, and databases

Hands-On



Install PostgreSQL & pgAdmin



Connect to database, explore catalogs

Assignment

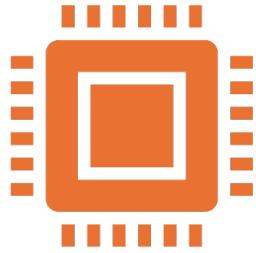


Install and configure PostgreSQL



Share screenshot of connected pgAdmin

Agenda



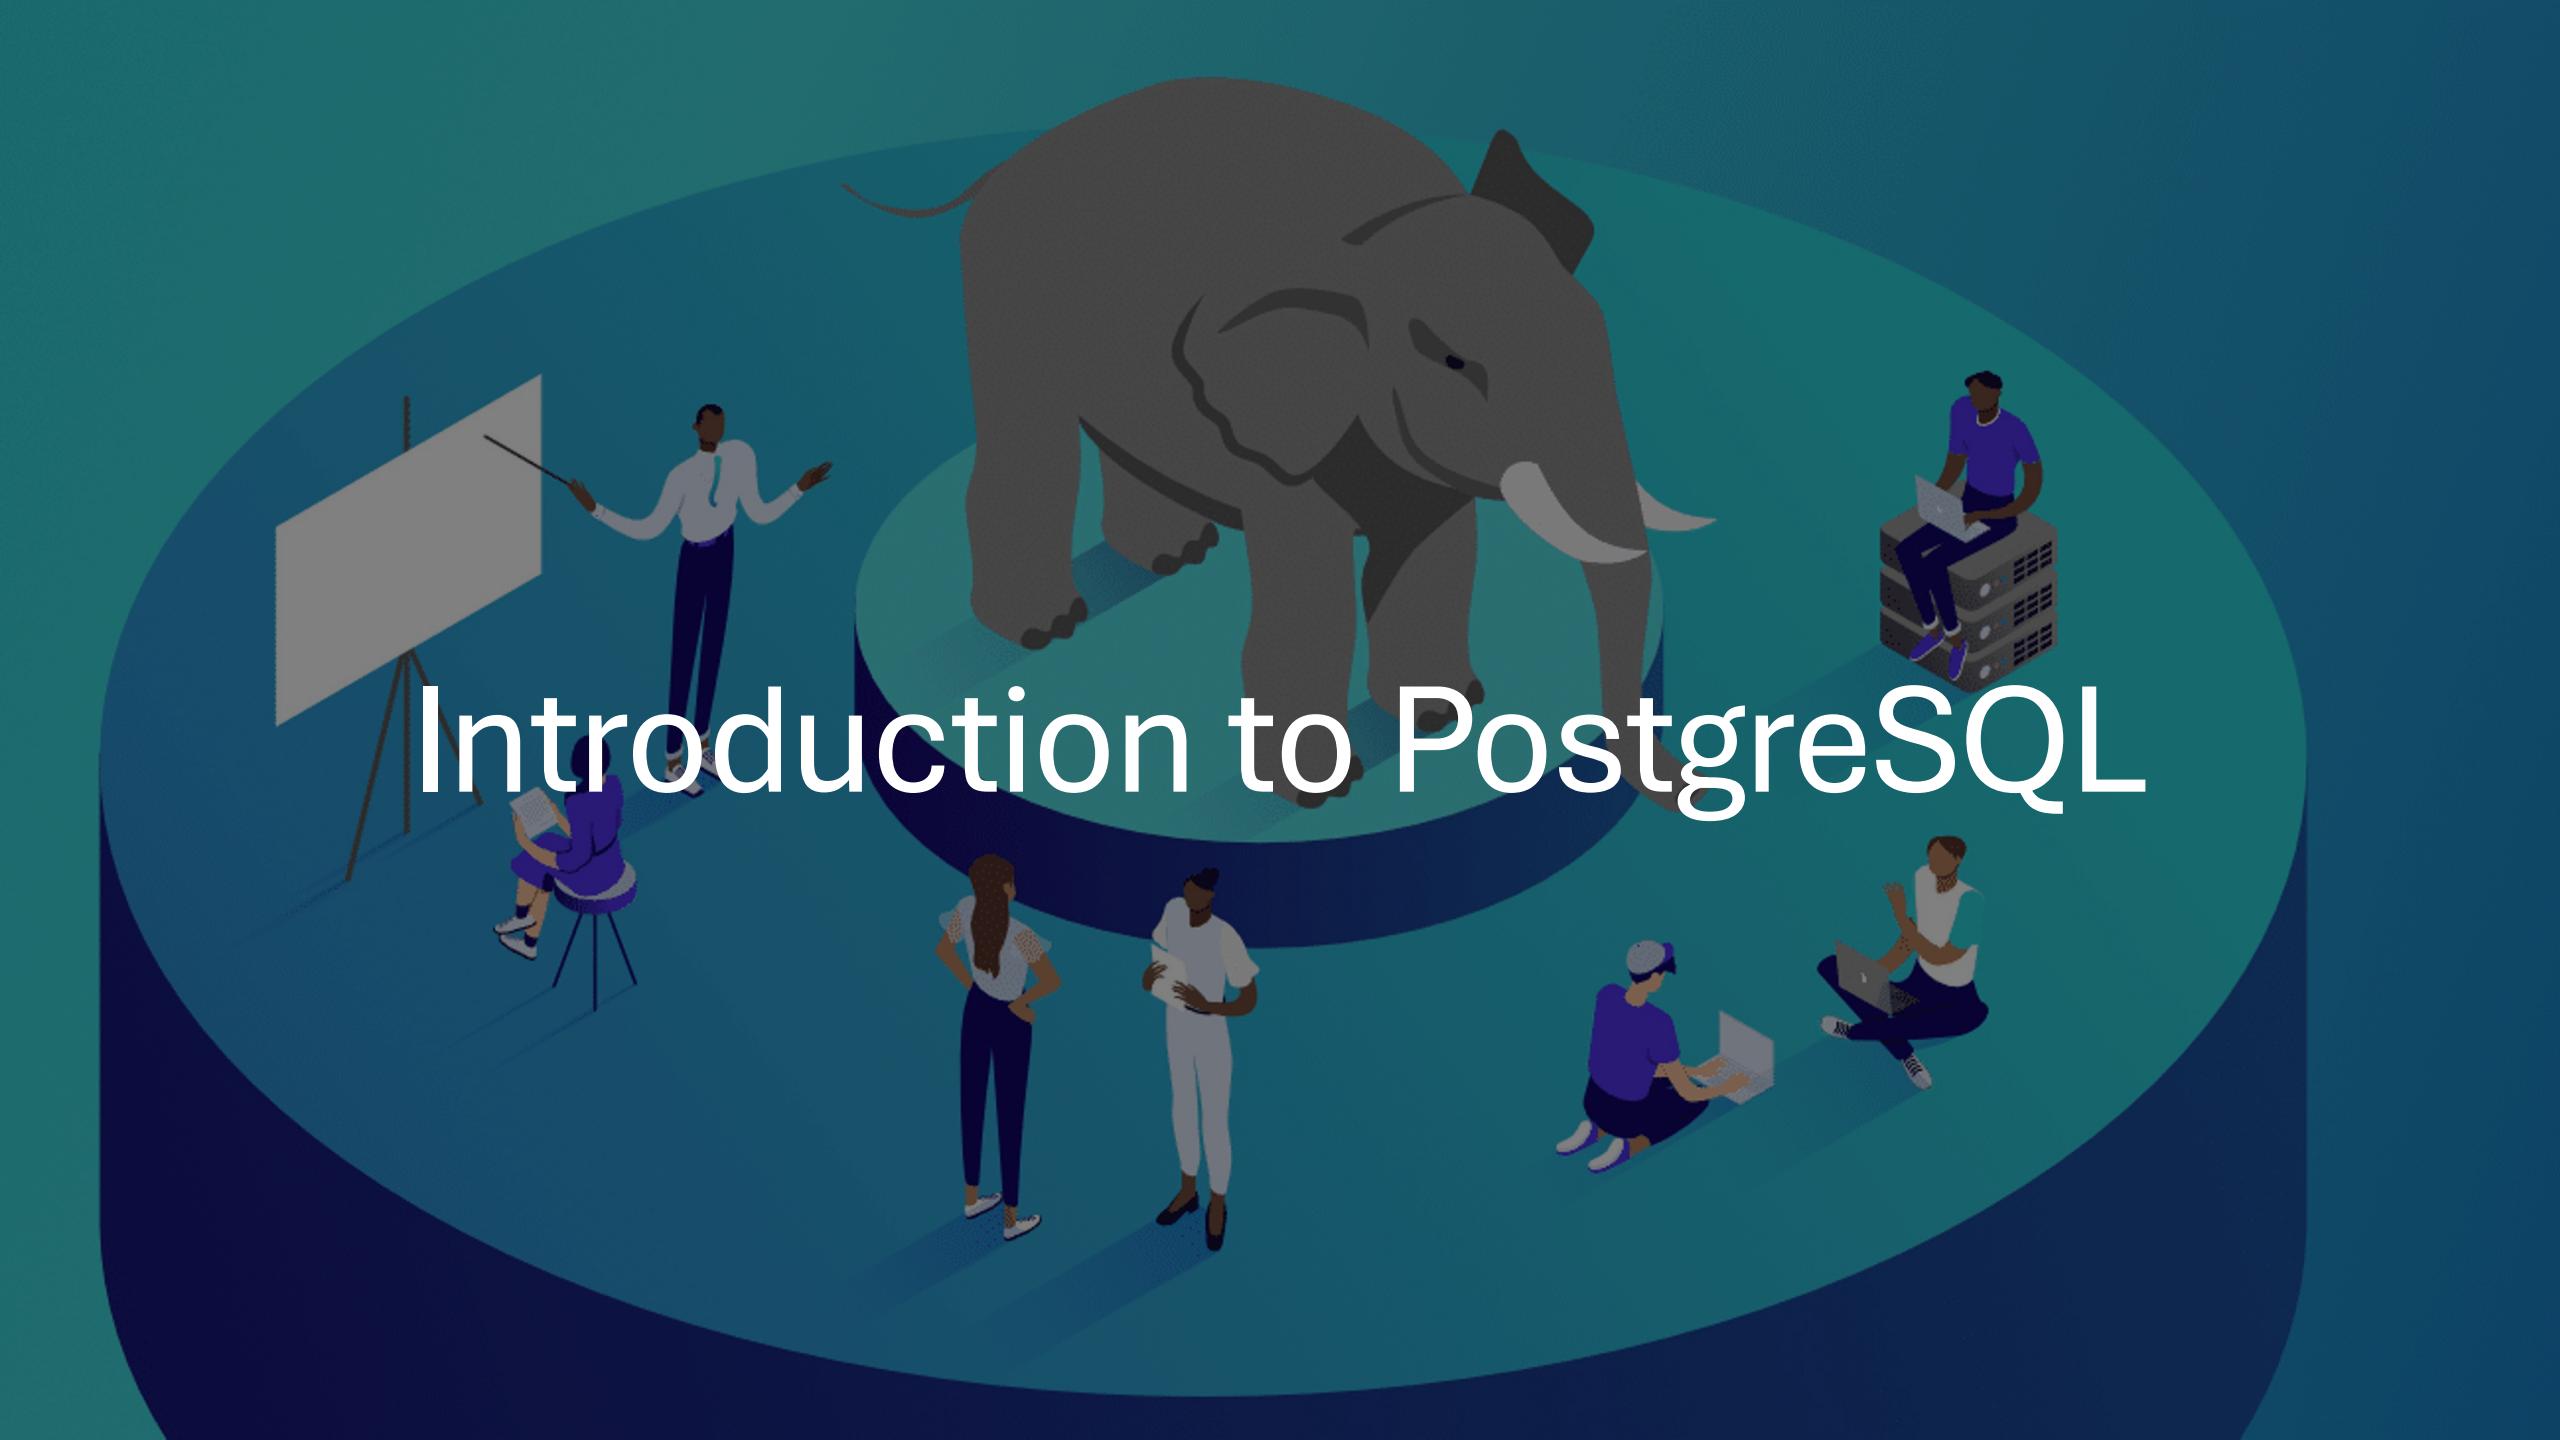
Installation (Windows,
macOS, Linux)



Configuration basics



Connecting to the
database



Introduction to PostgreSQL

What is PostgreSQL?



PostgreSQL (Postgres)



An advanced open-source



Relational Database Management System (RDBMS).

What is PostgreSQL?

Known for

Robustness

Extensibility

Standards compliance

What is PostgreSQL?



Designed to handle



Single-machine
applications



Large-scale data
warehousing



Web services

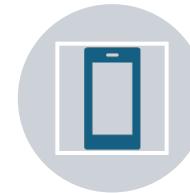
What is PostgreSQL?



Used for



Web



Mobile

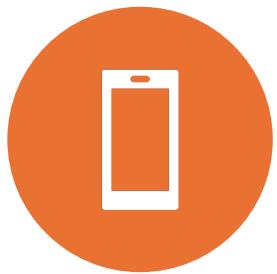


Geospatial



Analytics
applications

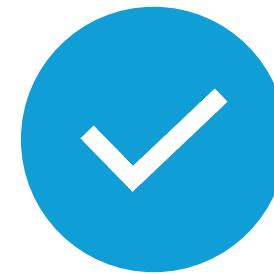
PostgreSQL 16



VARIETY OF
NEW FEATURES



IMPROVE
PERFORMANCE



USABILITY



SECURITY

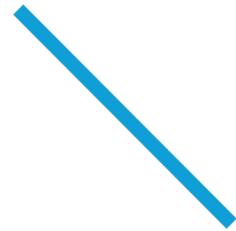
Performance Improvements



Query Performance
Enhancements



Partitioning
Improvements



Parallelism
Enhancements

Usability Enhancements

JSON and
JSONB
Enhancements

Enhanced
Window
Functions

Improved psql
Command-
Line Tool

Security Enhancements



Enhanced
Authentication and
Authorization



Advanced Data
Encryption

Administrative Features



Enhanced Backup and Restore



Improved Monitoring and Logging

Developer Features



Enhanced PL/pgSQL



Advanced Indexing



Key Contributors

- **Michael Stonebraker**
- The primary Architect
- Original Postgres project
- Key figure in database research

PostgreSQL Global Development Group (PGDG)



A group of
volunteers and
Companies



Oversee the
development



Maintenance of
PostgreSQL

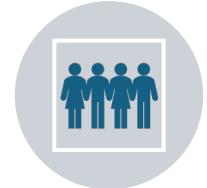
Community



Community-driven



Contributions from
developers
worldwide



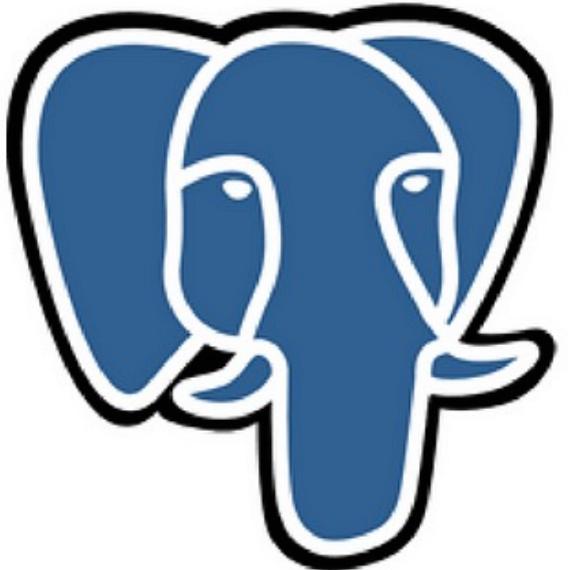
Known for its
collaborative spirit



Extensive
documentation



Active mailing lists



PostgreSQL

Design Philosophy

Standards Compliance



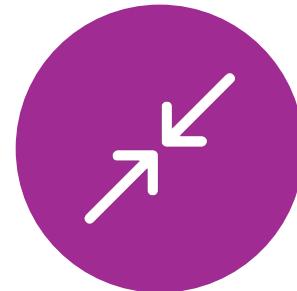
PostgreSQL adheres



To SQL standards



Ensuring compatibility



Ease of use

Extensibility

Designed to
be highly
extensible

Allowing
users to
define

New data
types

Functions

Operators

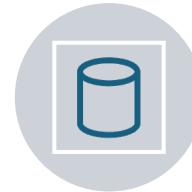
Reliability and Stability



Focuses on delivering



Stable



Reliable database system



Suitable for



Mission-critical applications

Open Source

Released under
the PostgreSQL
License

An open-source
license

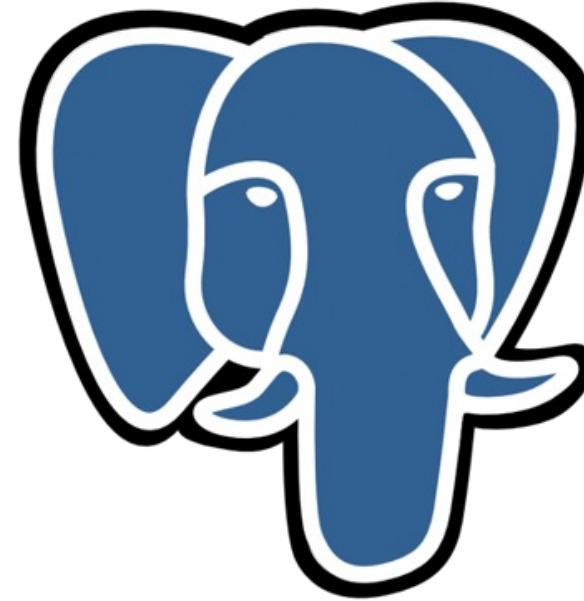
Allows for free
use

Modification

Distribution



Major Milestones



PostgreSQL

MVCC (1997)

Multi-Version Concurrency Control

A crucial feature

for high Concurrency

Performance

SQL Compliance (1996)



Transition from



PostQUEL to SQL



Aligning with



Industry
standards

Windows Support (2005)

Native support for

Windows

Broadening

PostgreSQL's user base

Streaming Replication (2010)



Hot standby



Improving



Data



Redundancy



Availability

Parallel Query Execution (2016)

Support for

Parallel
Query
Execution

Enhancing
performance

for large
queries.

Logical Replication (2017)



Introduction of



Logical
replication,



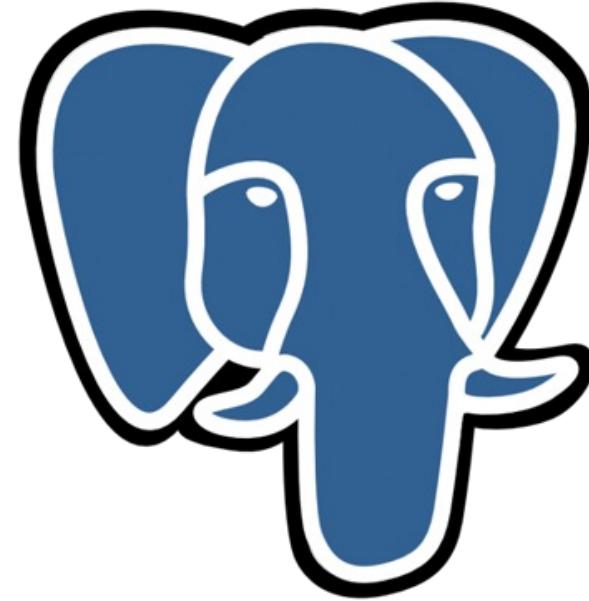
Allowing
selective



Replication of
data

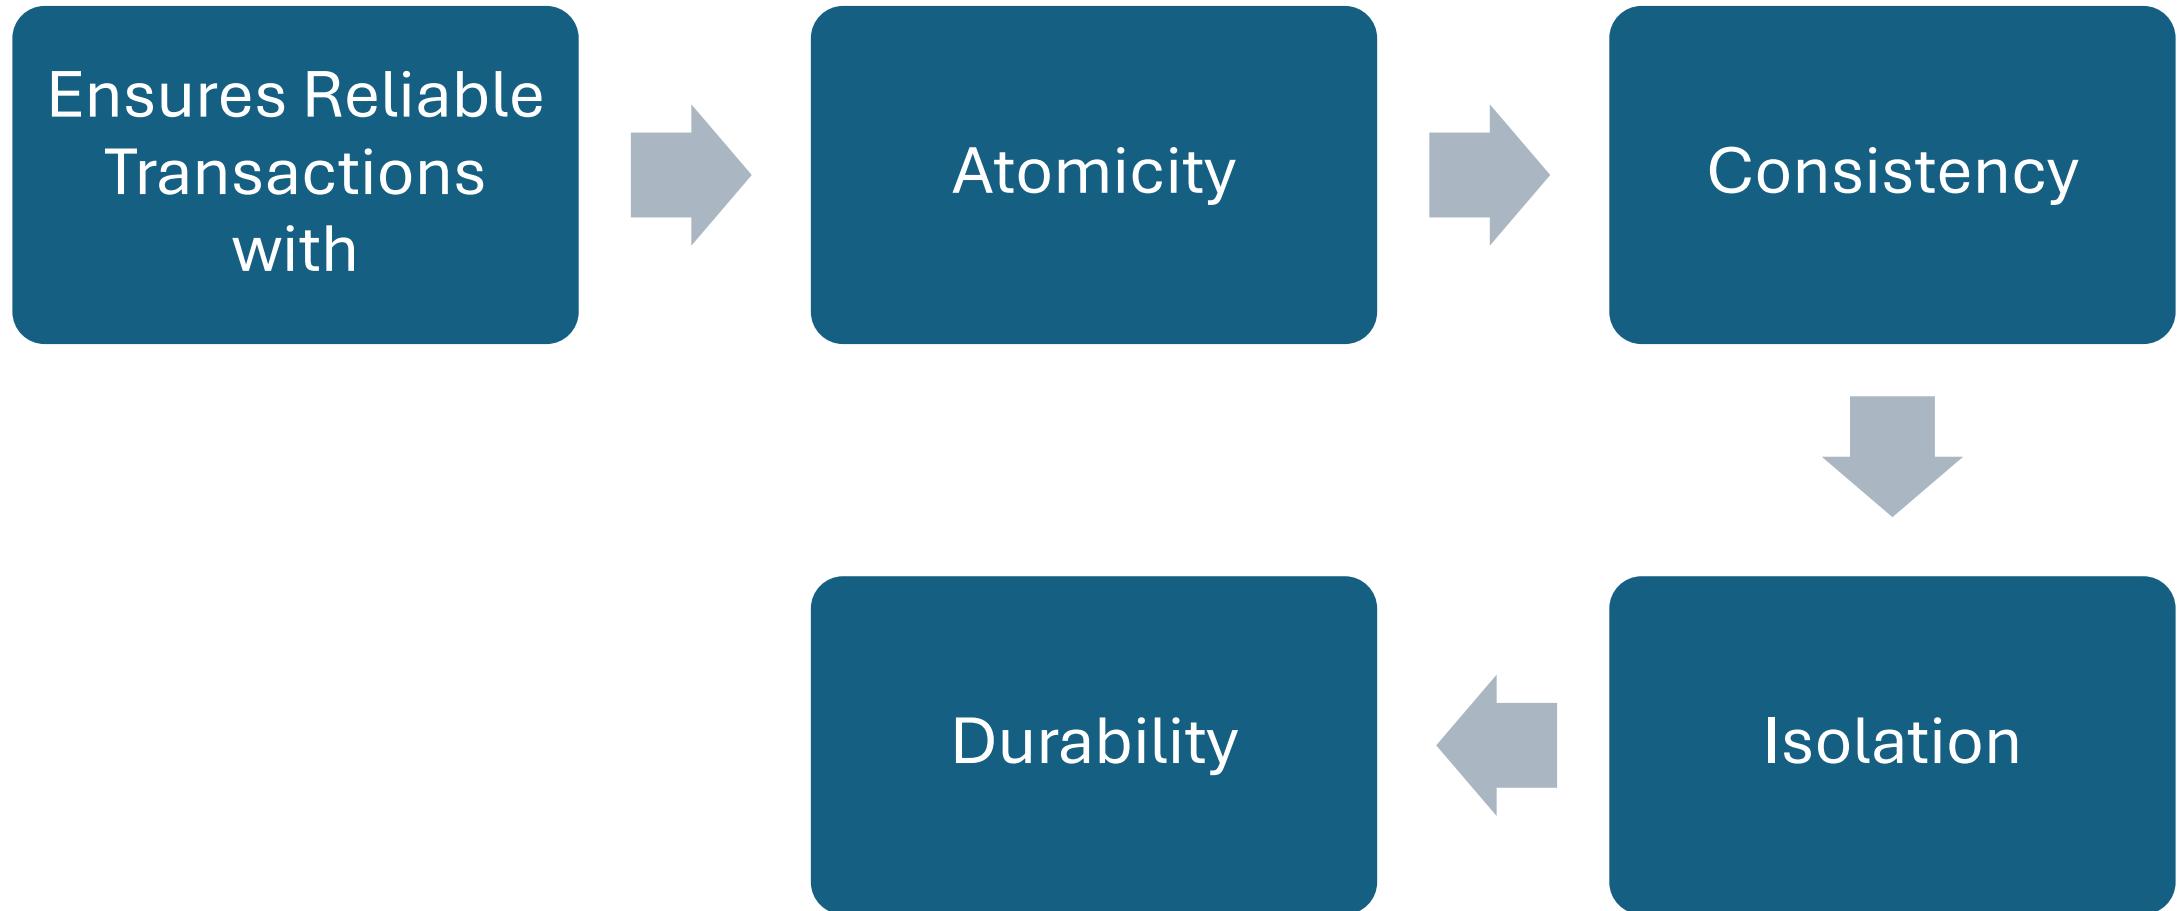


Key features and Benefits



Postgre^{SQL}

ACID Compliance



Atomicity

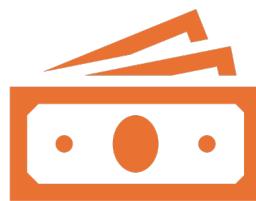
Each transaction is treated

Single "Atomic" Unit

Either completely succeeds or

Completely fails.

Atomicity



If any part of the transaction fails



Entire transaction is rolled back



Database remains unchanged.

Example

If a
transaction
involves

transferring
money from

Account A to
Account B.

Example



Either both the
debit from



Account A and



The credit to
Account B



Happen or
neither happens.

Example

If the debit succeeds

but the credit fails,

the transaction is rolled back.

Account A's balance is restored

to its original state

ACID Compliance



Provides robust
support



For concurrent
transactions



Without compromising
data integrity

Consistency

Ensures that a transaction

From one valid state

To another valid state,

Maintain all predefined rules

Constraints, and triggers

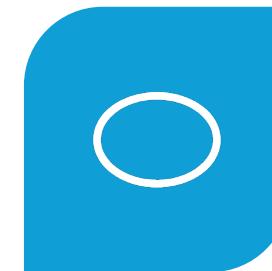
Consistency



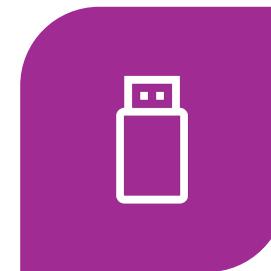
DATABASE
MUST



REMAIN
CONSISTENT

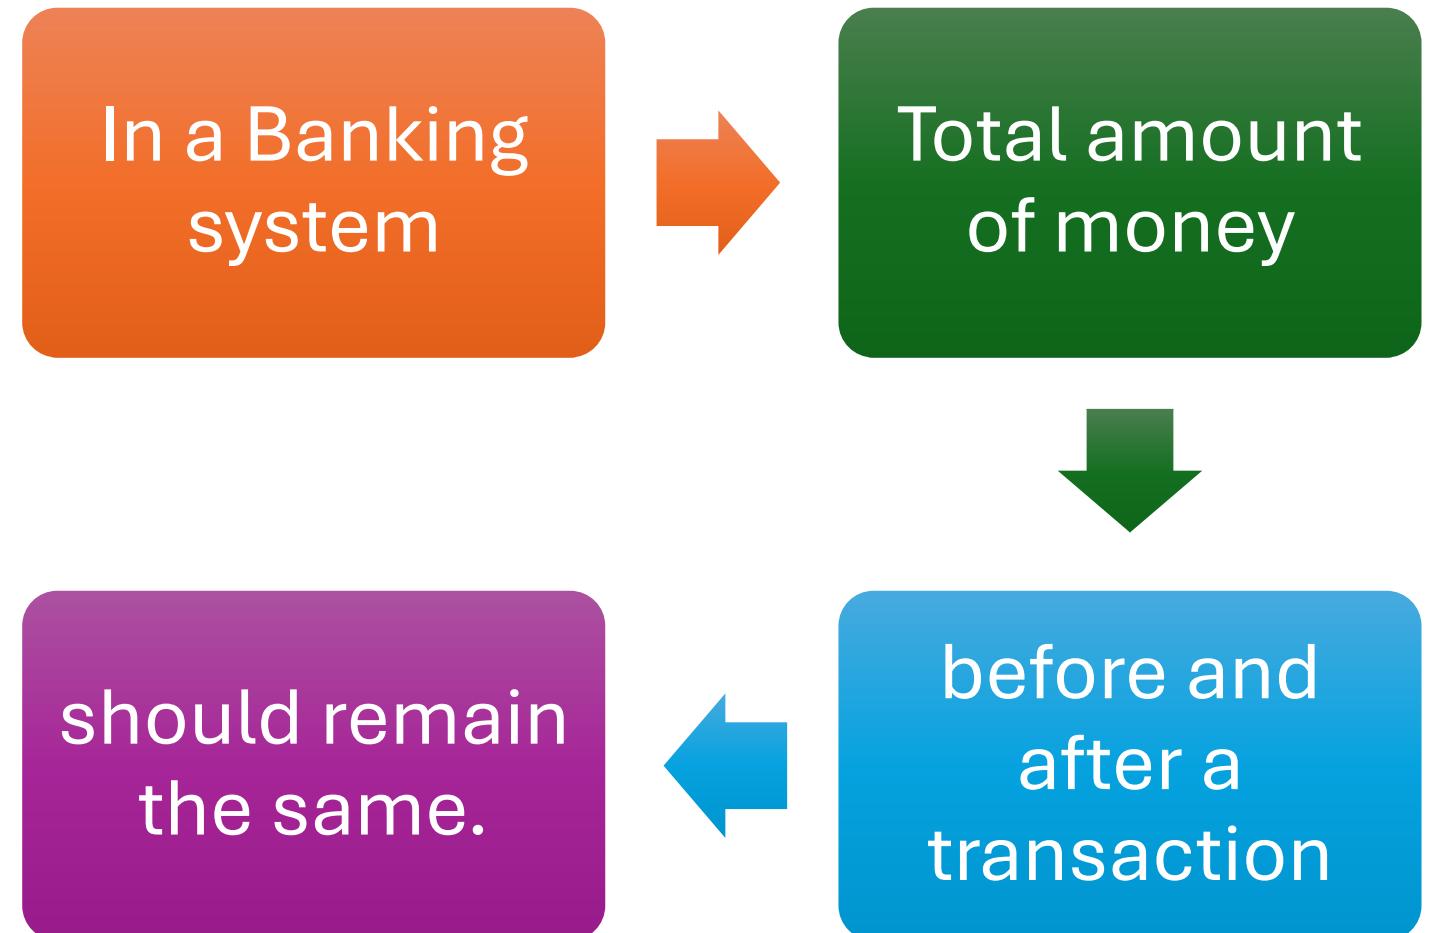


BEFORE AND
AFTER



TRANSACTION

Consistency



Isolation

Transactions

Executed in

Isolation

From One another

Isolation

If two transactions

Running Concurrently

One transferring money from

Account A to Account B

Another from Account C to Account D

Isolation

Isolation ensures

Transactions

Do not see each other's

Intermediate states

Durability

Once a transaction

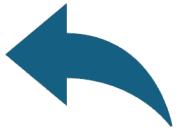
Has been committed

Will remain

Even in the event of

a system failure

Durability



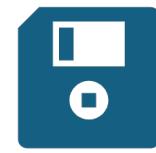
If a transaction
to update



A bank balance
is committed,



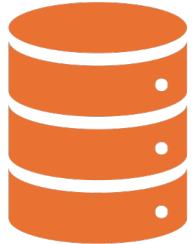
New balance is
guaranteed



to be saved in
the database.

Advanced Data Types

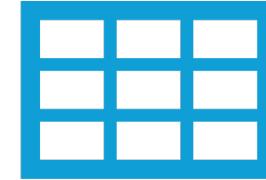
JSON/JSONB



Efficient storage



Querying of



JSON data.

Arrays



Support for



Multi-dimensional



Arrays

Hstore



Key-value pairs

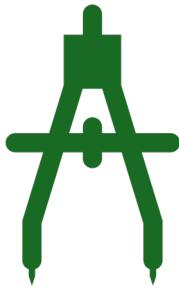


Storage

Geometric Types



Points,

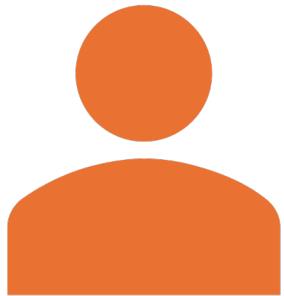


lines

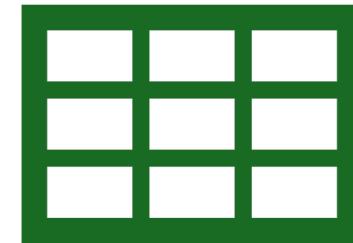


polygons

Custom Types



Users can define

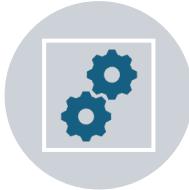


Own data types.

Extensibility



Users can
create



Custom
functions



Operators,

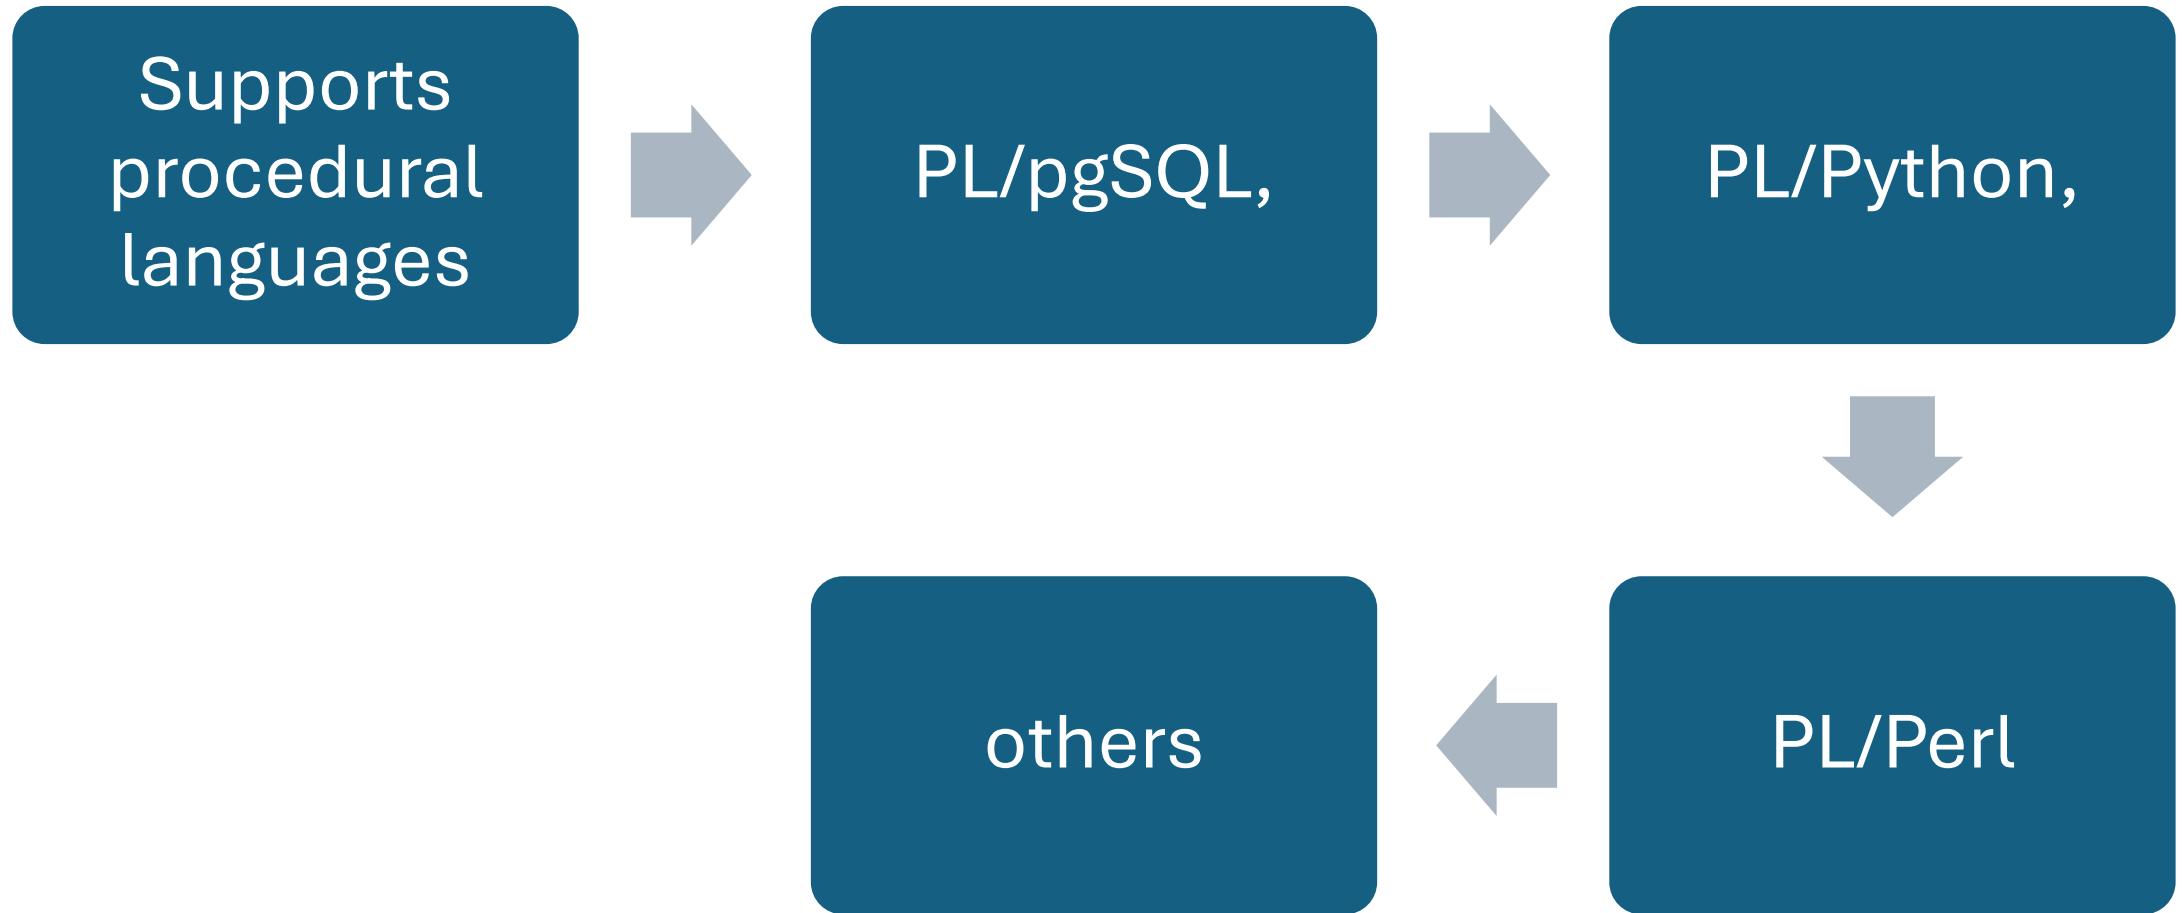


Data types



index types.

Extensibility



Extensibility



Extensions like



PostGIS (for
geospatial data)

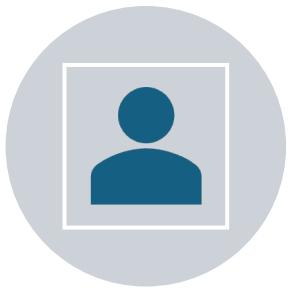


Full-text search



Easily integrated

Multi-Version Concurrency Control (MVCC)



Allows multiple transactions



to occur simultaneously without locking.



Ensures data consistency



High performance

Indexing



Supports various
index types



B-tree,



Hash,



GiST,

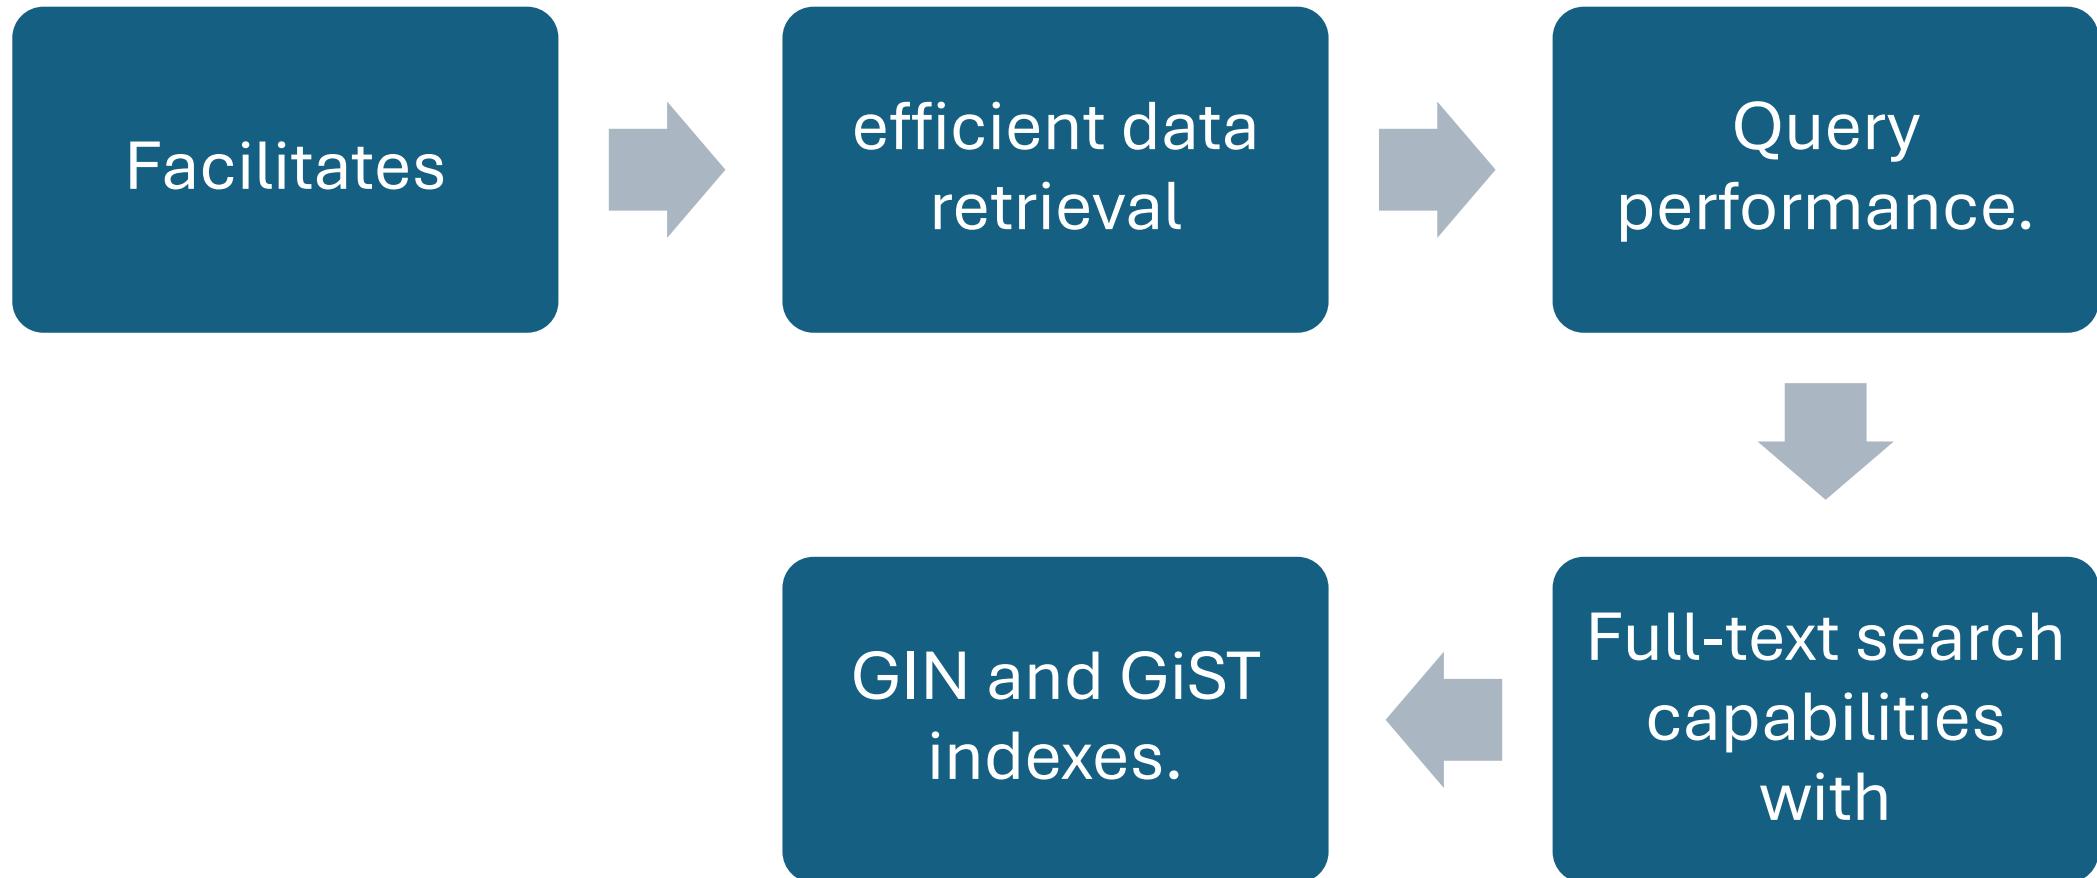


GIN



BRIN

Indexing



Foreign Data Wrappers (FDW)

Allows integration with

other databases and data sources.

Enables querying external data

as if it were a part of

the PostgreSQL database.

Full-Text Search



Built-in support for



full-text search
capabilities.



Indexing and querying
of



text data for fast search
results

Replication and High Availability

Supports both

Synchronous

Asynchronous

Replication

Replication and High Availability



STREAMING
REPLICATION



REAL-TIME DATA
REDUNDANCY



LOGICAL
REPLICATION



SELECTIVE DATA
SYNCHRONIZATION

Partitioning



Declarative
partitioning



for efficient
management



of large tables.

Partitioning

Supports

Range,

List,

Hash

composite partitioning methods.

Security



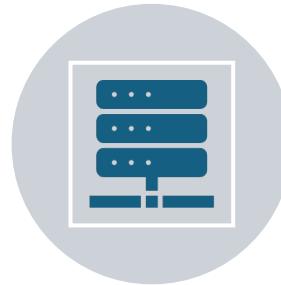
Robust access control
mechanisms



with roles and
permissions.



SSL support for

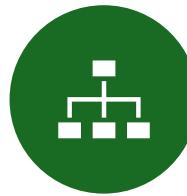


secure client-server
communication.

Security



Advanced
features like



row-level
security and



data
encryption.

Performance Optimization

Parallel Query
execution

for faster
processing

of large
datasets.

Performance Optimization

Query planner

Optimizer

for efficient

Query execution

Performance Optimization



CACHING MECHANISMS



TO IMPROVE
PERFORMANCE.

Backup and Recovery

Comprehensive backup and

restore options with

tools like pg_dump and

pg_basebackup.

Backup and Recovery



Point-in-time
recovery (PITR)



for restoring
databases



to a specific state.

Benefits of PostgreSQL

Open Source and Free

PostgreSQL is open source,
allowing for free usage,
modification, and distribution.
No licensing costs,
making it cost-effective
for organizations

Flexibility

Highly extensible and
adaptable to various use cases.
Suitable for a wide range of applications,
from small-scale to
large-scale enterprise systems.

Community and Support



Strong, active community



providing extensive documentation,



support, and regular updates.



Numerous third-party tools and



extensions available.

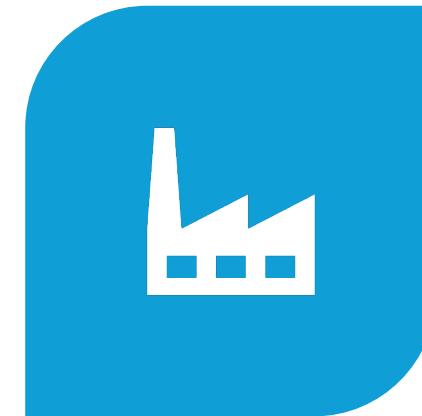
Reliability and Stability



PROVEN TRACK
RECORD OF



RELIABILITY AND
STABILITY IN



PRODUCTION
ENVIRONMENTS.

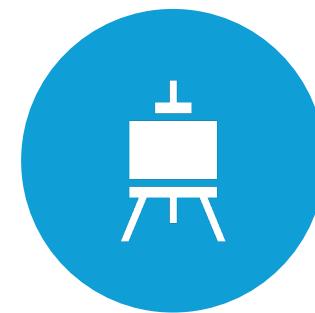
Standards Compliance



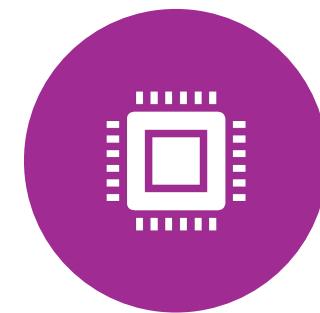
ADHERES TO SQL
STANDARDS,



ENSURING
COMPATIBILITY AND



EASE OF USE FOR
DEVELOPERS



FAMILIAR WITH
SQL.

Standards Compliance



REGULAR UPDATES



TO MAINTAIN
COMPLIANCE



WITH EVOLVING
STANDARDS.

Cross-Platform Compatibility



Runs on various operating systems



including Windows, macOS, and Linux.



Supports integration with



Various programming

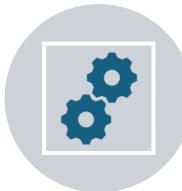


languages and frameworks.

Performance



Advanced indexing and



query optimization
techniques



ensure high
performance.



Scalability to
handle



large volumes of
data and



high transaction
loads.

Data Integrity

Strong support for data integrity

through constraints,

triggers, and foreign keys.

Ensures accurate and

consistent data storage and

retrieval.

Geospatial Capabilities

PostGIS extension

provides robust support

for geographic information systems (GIS).

Ideal for applications

requiring spatial data management.

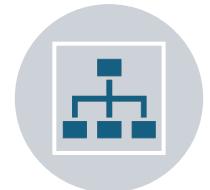
Enterprise Features



Comprehensive
feature set



that meets the
needs of



enterprise-level
applications.



Support for
complex queries,



transactions



data analytics.

Agenda



PostgreSQL vs Oracle: Key architectural differences



PostgreSQL installation & tools (pgAdmin, DBeaver)



Overview of pg_catalog, schemas, roles, and databases

Hands-On



Install PostgreSQL & pgAdmin

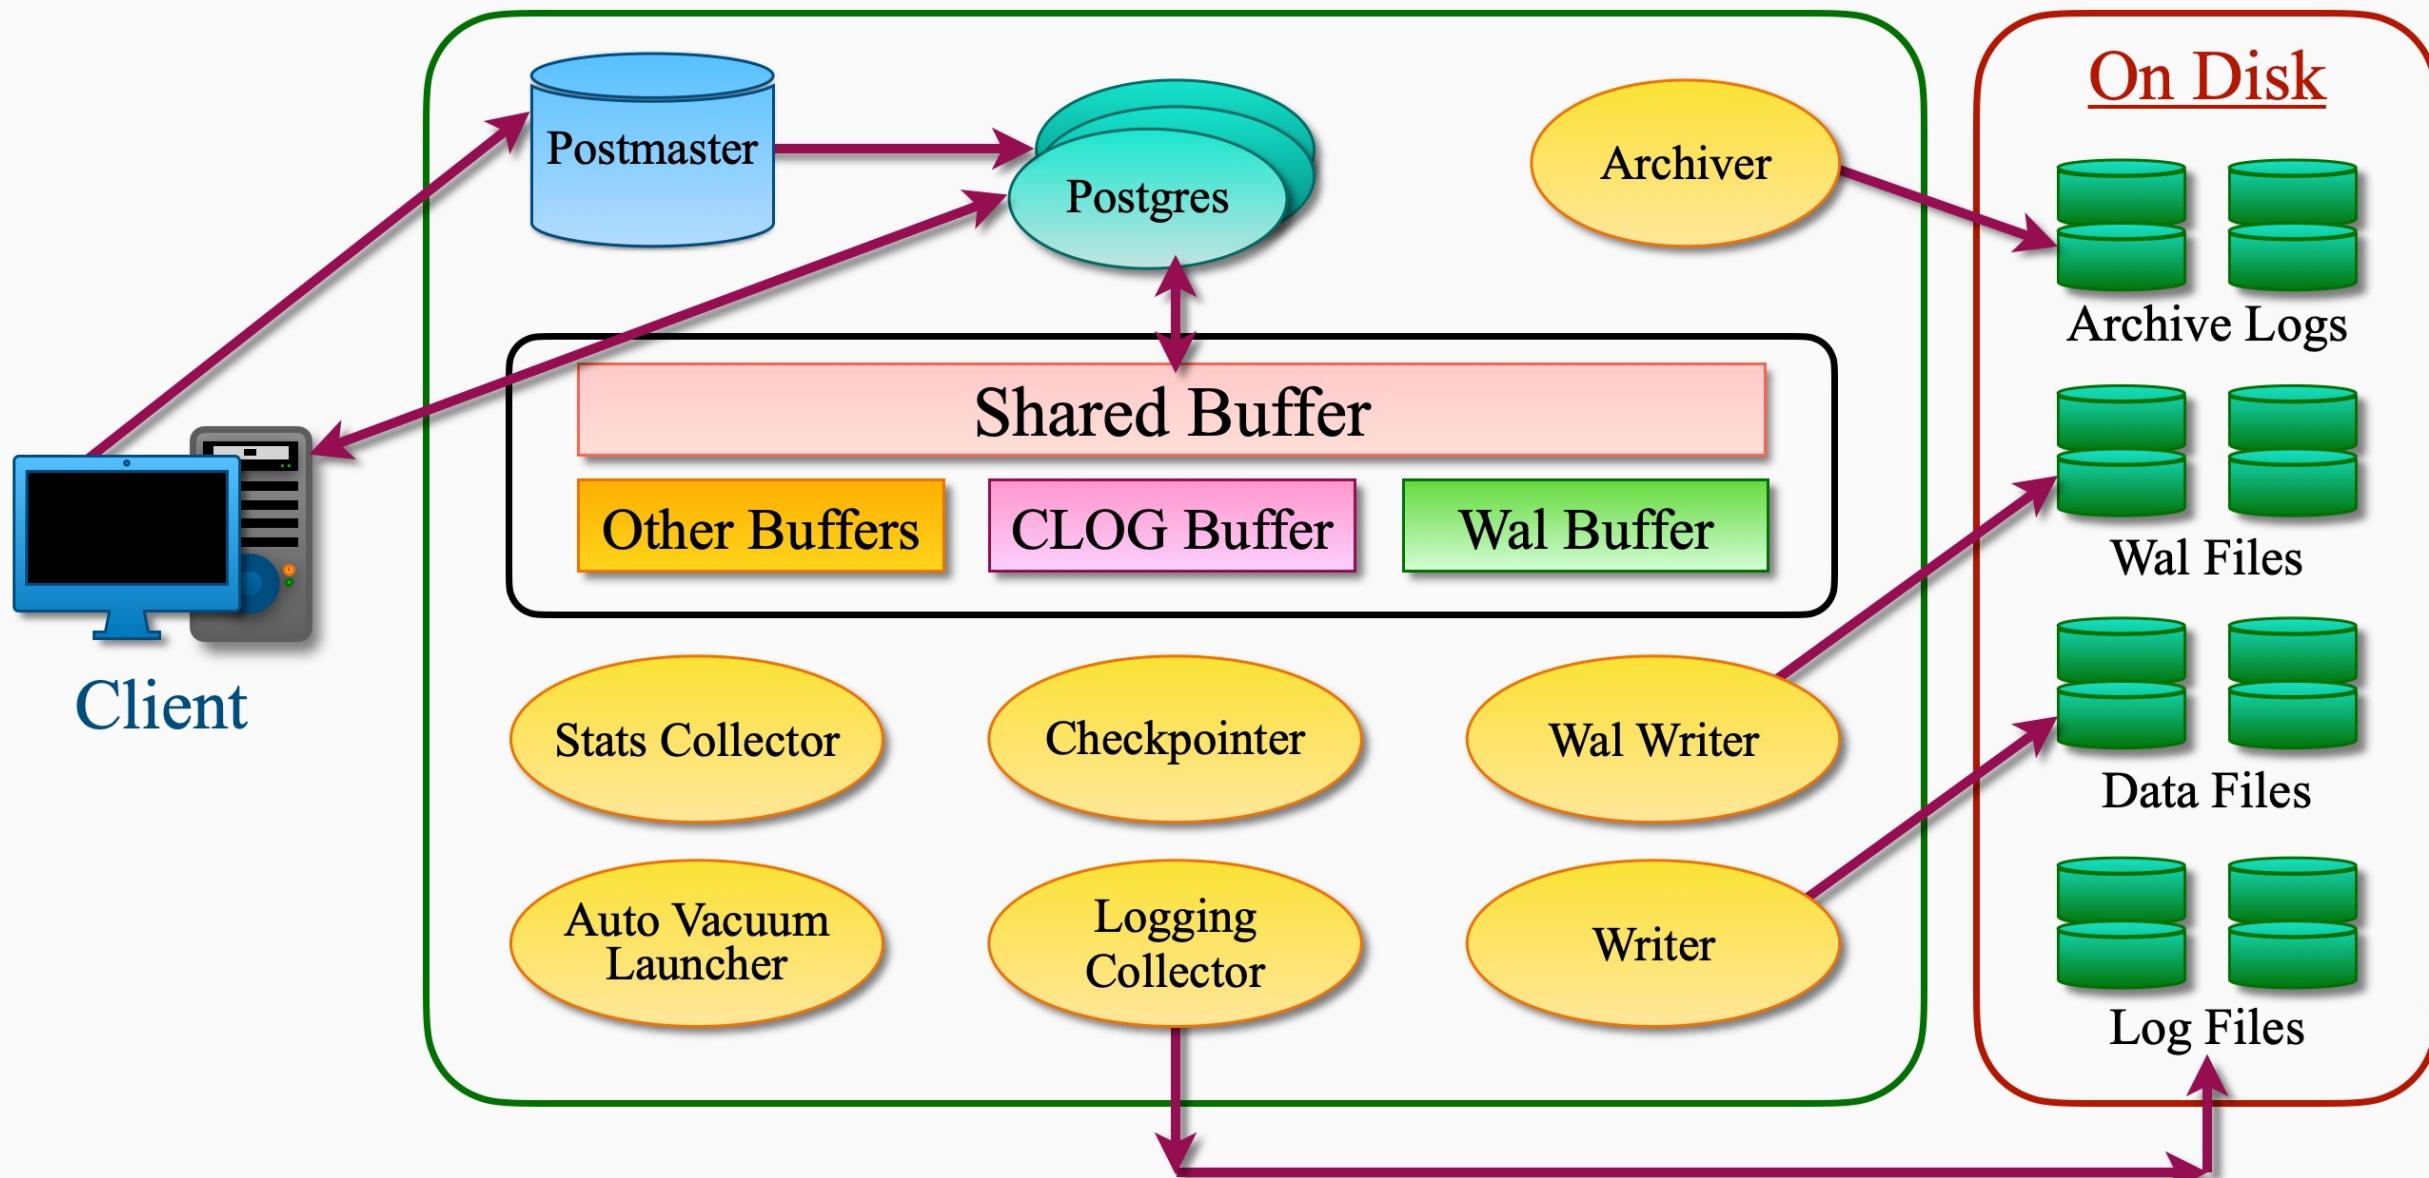


Connect to database, explore catalogs

PostgreSQL vs Oracle: Key architectural differences



PostgreSQL Process and Memory Architecture



PostgreSQL Architecture Key components

1. Postmaster Process (Supervisor)

2. Shared Memory Segments

3. Background Processes

4. Physical File Structure

1. Postmaster Process (Supervisor)

The **first process**

Started when PostgreSQL launches

1. Postmaster Process (Supervisor)

Acts as a **listener**, handles:

Authentication

Authorization

Spawning new backend processes

Monitoring and restarting failed processes

2. Shared Memory Segments

Memory areas used for

query execution,

data caching, and

maintenance tasks.

Shared Buffers



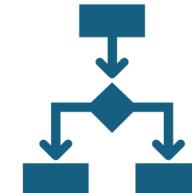
Used for all **read/write operations**



Modified (dirty) pages
are written to disk



by the **background writer**



Controlled by:
`shared_buffers`

WAL Buffers (Write-Ahead Logging)

Temporarily store metadata of changes

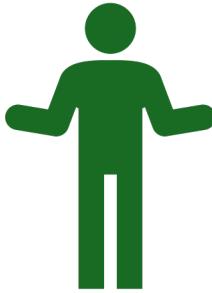
Ensures durability and crash recovery

Controlled by: wal_buffers

CLOG Buffers (Commit Logs)



Track **transaction status**



(committed, aborted,
etc.)



Shared across all
sessions

Work Memory

Used for sorting, hashing, joins

Controlled by: `work_mem`

Separate allocation for each operation

Maintenance Work Memory

Used for **index creation,**

foreign key additions, etc.

Controlled by:

`maintenance_work_mem`

Temp Buffers

Used by **temporary tables**

Session-specific

3. Background Processes



Ensure



data consistency



help



maintain database
health.

Background Writer



Writes



dirty pages



from shared
buffers



to disk

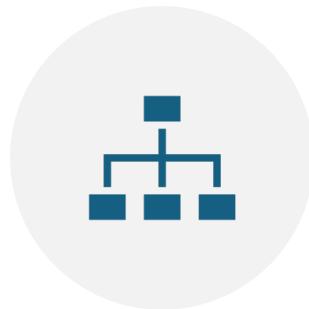
Background Writer



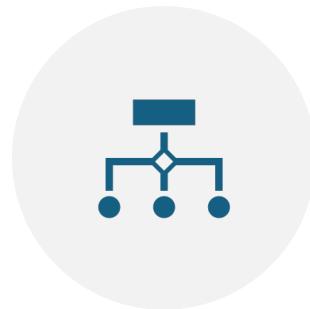
CONTROLS:



BGWRITER_DELAY



BGWRITER_LRU_MAXPAGES



BGWRITER_LRU_MULTIPLIER

Checkpointer

Periodically **flushes**

dirty pages to disk

Essential for

crash recovery

Checkpointer

Triggered by:

Timeout (checkpoint_timeout)

Manual CHECKPOINT command

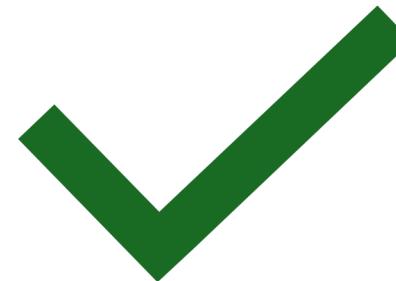
Backup processes

max_wal_size limit

WAL Writer



Writes data from **WAL buffer** to
WAL files on disk



Ensures write operations are
persistent before commit

Autovacuum Launcher



Automates **VACUUM** and
ANALYZE



Maintains **table statistics** and
bloat control

Statistics Collector

Gathers usage metrics:

Query counts

Table scans

Function usage

Statistics Collector

Controlled by:

track_activities

track_functions

track_counts

track_io_timing

Logging Collector

Captures server logs to files

Controlled by:

`logging_collector = ON`

`log_directory = 'pg_log'`

Archiver

Saves WAL segments to archive storage

for point-in-time recovery

Enabled during archive mode

4. Physical File Structure



PostgreSQL uses several file types



to store and manage data persistently:

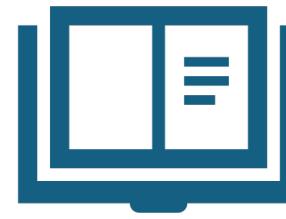
Data Files



Store **actual user data**



Accessed via shared
buffers



for reading/writing

WAL Files



Store **change logs** (Write-Ahead Logging)



Used for crash recovery

Log Files

Store **server logs** including

errors,

warnings, and

execution info

Archive Logs

Archived WAL
segments

for recovery
purposes

Workflow Summary

Client sends a query → received by Postmaster

Postmaster forks a backend (Postgres) process

Data retrieved from Shared Buffers

(or disk if not cached)

Workflow Summary

If write operation:

Change goes to Shared Buffer → WAL Buffer

WAL Writer commits to WAL file

Background Writer/Checkpointer writes to disk

Statistics collected → Logging → Archiving if enabled

Conclusion



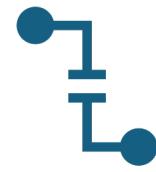
PostgreSQL's architecture is



robust, scalable,
and



designed for high
reliability,

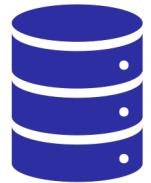


fault-tolerance,
and concurrency.

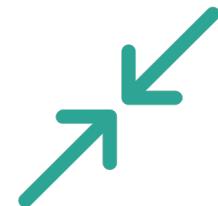
Conclusion



Its modular background
processes and



memory components

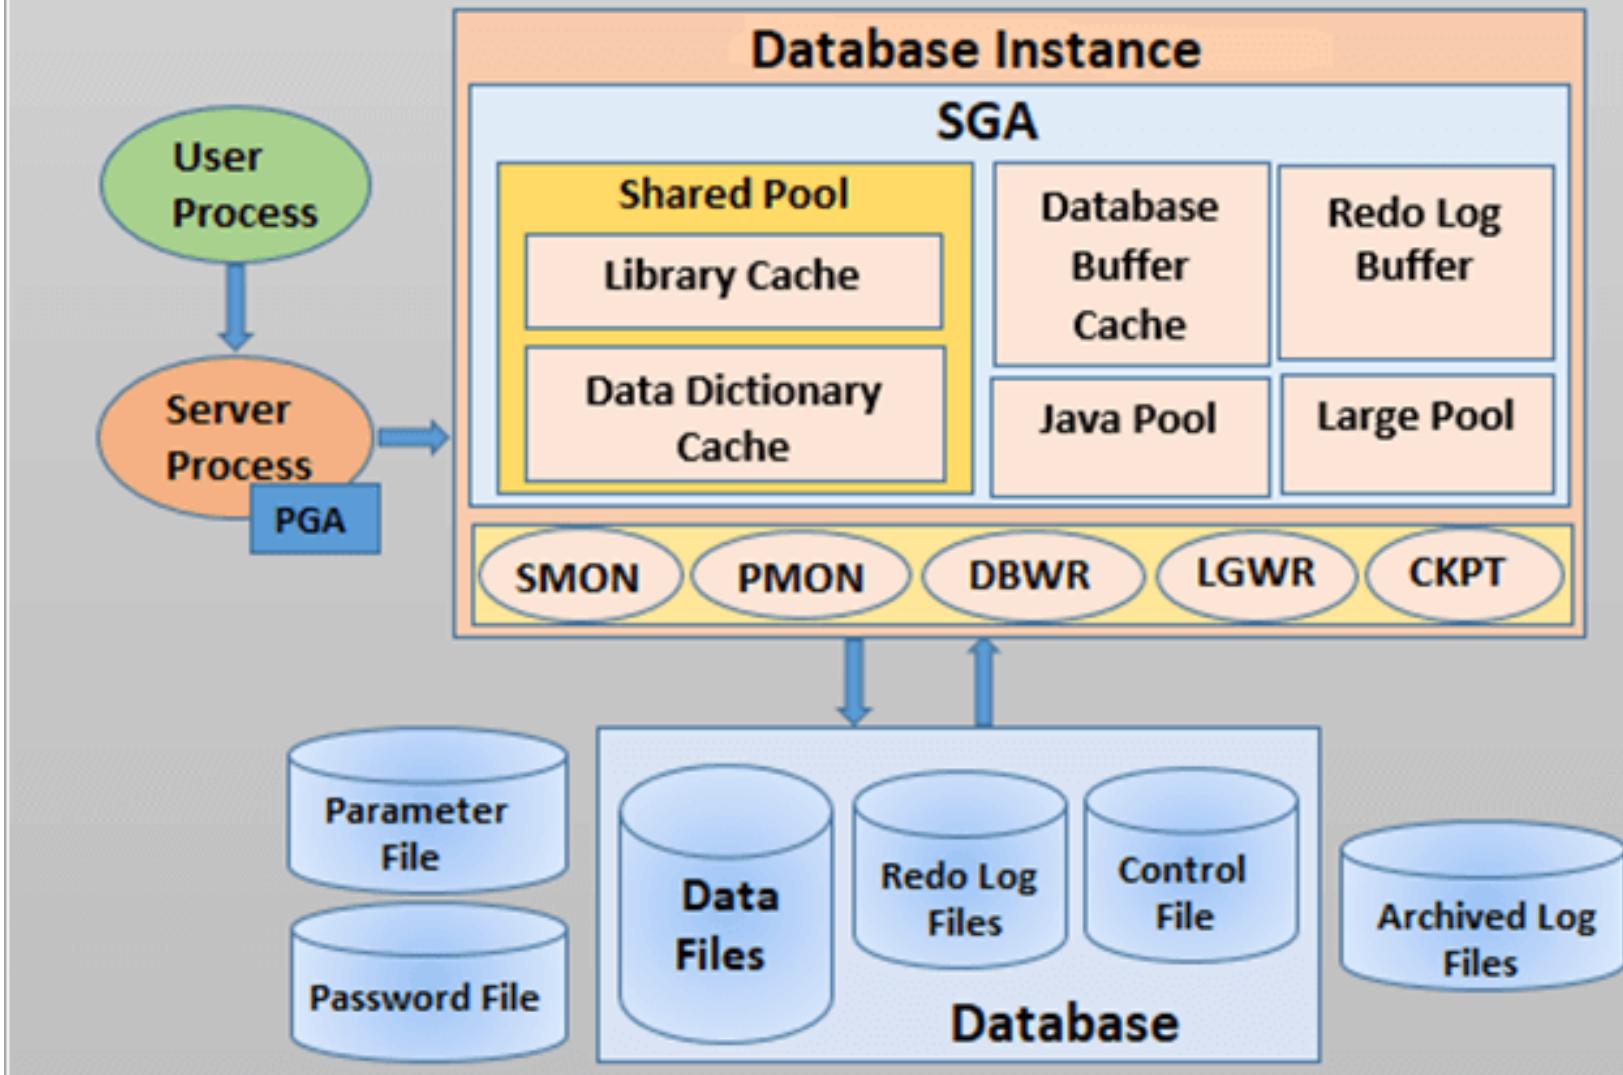


ensure smooth
operations and quick
recovery,



making it a strong choice
for modern applications.

Oracle Database Architecture



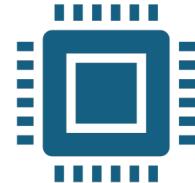
Oracle Database Architecture



Oracle Database is a
multi-model



**relational database
management system**



known for its scalability,
robustness, and



enterprise-grade
features.

Oracle Database Architecture



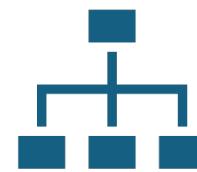
Designed to manage



large amounts of data
efficiently



using a combination of
memory,



processes, files, and
logical storage
structures.

Oracle Database Key Components

Oracle Instance (Memory + Processes)

Physical Storage Structures

Processes

Logical Storage Structures

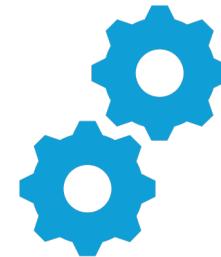
1. Oracle Instance (Memory + Processes)



Combination of:



Memory Structures
(SGA & PGA)

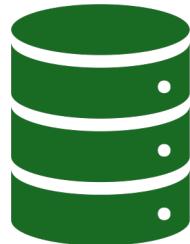


Background Processes

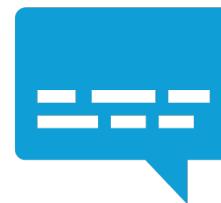
1. Oracle Instance (Memory + Processes)



The instance manages



data stored in physical
files and



serves client requests.

1.1 Memory Structures

A. SGA (System Global Area) – Shared memory

Allocated when the instance starts,

shared among all server/background processes.

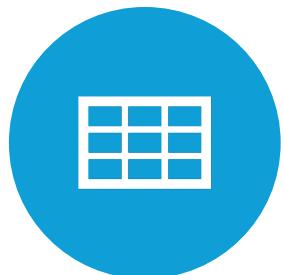
Key components



Database Buffer Cache:



Caches recently used data blocks.



Redo Log Buffer:



Stores redo entries (change vectors) for data recovery.

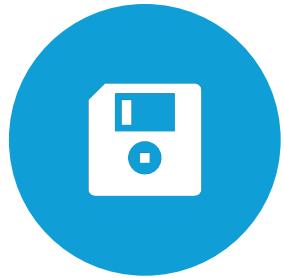
Key components



Shared Pool:



Caches SQL statements, PL/SQL code, and metadata.



Library Cache:



Part of Shared Pool;
stores parsed SQL and PL/SQL.

Key components

Data Dictionary Cache:

Stores definitions of tables, users, privileges.

Large Pool:

Used for large memory operations (e.g., parallel queries).

Key components

Java Pool: Used for Java-based stored procedures.

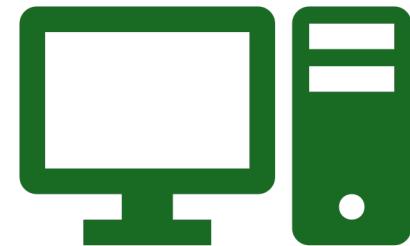
Controlled by SGA_MAX_SIZE and

init.ora/spfile.ora parameters.

B. PGA (Program Global Area)



Private memory

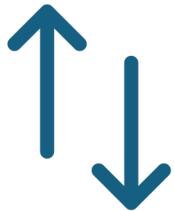


Dedicated to each server
process

B. PGA (Program Global Area)



Stores session-specific data like:



Sort areas



Cursor state



Session variables

Background Processes



Are essential
daemons



that help manage
memory,



I/O, logging, and



recovery

Mandatory Background Processes

SMON (System Monitor):

Performs crash recovery, temp space cleanup, etc.

PMON (Process Monitor):

Cleans up failed user processes and resources.

Mandatory Background Processes



DBWR (Database Writer):



Writes dirty buffers from memory to datafiles.



LGWR (Log Writer):



Writes redo log buffer to redo log files.

Mandatory Background Processes



CKPT (Checkpoint Process):



Signals DBWR and updates
control & data files.

Optional Background Processes



ARCn (Archiver):



Archives redo logs for recovery.



CJQ0 (Job Queue Coordinator):

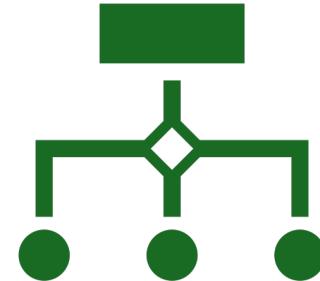


Manages scheduled jobs.

Optional Background Processes



SMCO (Space Management Coordinator):



Handles tasks like shrinking segments.

2. Physical Storage Structures

File Type	Description
Data Files	Store actual table/index data.
Redo Log Files	Store all changes made to the DB, used for recovery.
Control Files	Contain metadata about the database structure and state.
Archived Redo Logs	Backups of redo logs used for point-in-time recovery.
Parameter File (PFILE/SPFILE)	Configuration file for instance startup.
Password File	Stores authentication info for privileged users.

3. Processes



Processes manage



interaction between



users and the database
engine.

3.1 User Process



Initiated when a user connects



Handles communication



between client and server

3.2 Server Process

Executes SQL
queries

on behalf of the
user process

Reads/writes
data to memory
or files

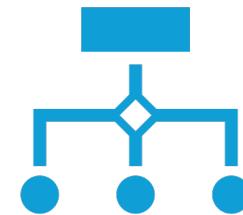
3.3 Background Processes



Run automatically



to handle tasks like
recovery,



buffer management,
etc.

4. Logical Storage Structures

Structure	Description
Tablespace	Logical storage unit composed of one or more data files
Segment	Set of extents used to store a specific database object
Extent	Group of contiguous data blocks
Data Block	Smallest unit of storage, defined by DB_BLOCK_SIZE (e.g., 8KB)

Oracle Workflow (Simplified)

User connects → User & Server processes are created

Query Execution:

SQL parsed and cached (Shared Pool)

Data fetched from **Buffer Cache** (or disk if not cached)

Query Execution

Modifications → written to **Buffer Cache and Redo Log Buffer**

LGWR writes Redo Buffer to Redo Logs

DBWR writes changes to Data Files

(on checkpoint or buffer pressure)

Recovery (if needed):



Redo logs and



archived logs



used to restore DB

Summary Table

Component	Details
Instance	SGA + PGA + Background Processes
Memory Structures	SGA (shared), PGA (private)
Mandatory Processes	SMON, PMON, DBWR, LGWR, CKPT
Physical Files	Data Files, Redo Logs, Control Files
Logical Units	Tablespaces → Segments → Extents → Data Blocks

Why Compare PostgreSQL & Oracle?



Understanding architectural differences helps in:



Migration planning



Cost-benefit analysis



Technology stack decisions

License & Philosophy

Feature	PostgreSQL	Oracle
License	Open Source (PostgreSQL License - similar to BSD)	Commercial (Proprietary License)
Philosophy	Extensibility, Standards Compliance, Community-Driven	Feature-Rich, Enterprise-Grade, Vendor Lock-in

Architecture Overview

Aspect	PostgreSQL	Oracle Database
Process Model	Multi-Process (One process per client connection + background processes)	Multi-Threaded (Single process with multiple threads + background processes)
Memory Architecture	Shared Buffers, Work Mem, WAL Buffers	SGA (System Global Area), PGA (Program Global Area), Redo Log Buffer

Architecture Overview

Aspect	PostgreSQL	Oracle Database
Storage Engine	Single storage engine	Multiple engines (Heap Organized Tables, Index Organized Tables, etc.)
Extensibility	Highly extensible (user-defined types, operators, functions, custom languages)	Limited extensibility (supports PL/SQL, but not new data types easily)

Transaction & Concurrency Control

Feature	PostgreSQL	Oracle
MVCC (Multi-Version Concurrency Control)	Native MVCC (Row-versioning, no read locks for SELECTs)	MVCC via Undo Segments (Undo data in rollback segments)
Locking	Row-level locks, advisory locks	Row-level locks, fine-grained locking, table-level partition locks
Isolation Levels	Read Committed, Repeatable Read, Serializable	Read Committed, Serializable, Read-Only Transactions

Replication & High Availability

Feature	PostgreSQL	Oracle
Replication	Streaming Replication, Logical Replication (built-in), 3rd party tools (e.g., Bucardo)	Oracle Data Guard (physical), Oracle GoldenGate (logical)
Sharding	External tools (Citus, pg_shard)	Native (Oracle Sharding)

Data Types & SQL Compliance

Feature	PostgreSQL	Oracle
Data Types	Rich: JSON/JSONB, Arrays, HSTORE, Range Types	Limited: JSON support added in 12c but less flexible
Standards	Very close to ANSI SQL compliance	SQL compliance + proprietary PL/SQL

Partitioning

Feature	PostgreSQL	Oracle
Partitioning	Declarative partitioning (since PostgreSQL 10)	Advanced partitioning (list, range, hash, composite)
Management	Manual in earlier versions; improved automation in recent versions	Fully automated, advanced features

Backup & Recovery

Feature	PostgreSQL	Oracle
Backup	pg_dump (logical), base backups + WAL archiving	RMAN (block-level backup), Data Pump (logical export/import)
Point-in-Time Recovery	WAL-based PITR	Redo/Undo Logs, Flashback Database

PL/SQL vs PL/pgSQL

Feature	PostgreSQL	Oracle
Procedural Language	PL/pgSQL (open), also supports Python, Perl, etc.	PL/SQL (proprietary, tightly integrated)
Triggers/Functions	Easily create functions in multiple languages	Primarily in PL/SQL

Performance & Optimization

Feature	PostgreSQL	Oracle
Optimizer	Cost-based, extensible, adaptive plans in progress	Highly advanced Cost-based Optimizer, adaptive plans, hints
Parallel Query	Supported but less mature	Advanced, with fine-grained control

Security & Auditing

Feature	PostgreSQL	Oracle
Authentication	SSL, GSSAPI, LDAP, Kerberos	Advanced (including fine-grained auditing, TDE, DB Vault)
Auditing	Basic logging and extensions (e.g., pgaudit)	Comprehensive built-in auditing

When to Use What?

PostgreSQL	Oracle
Startups, Open Source Projects	Large Enterprises, Legacy Systems
Cost-sensitive environments	High compliance, mission-critical
Customization & Extensions	Advanced features, support

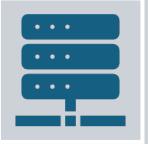
Summary of Key Differences

PostgreSQL	Oracle
Open-source, community-driven	Proprietary, enterprise-focused
Process-based architecture	Thread-based shared memory architecture
Extensible with custom data types	Fixed feature set, but rich enterprise features
MVCC via row versioning	MVCC via undo segments
Free, with optional paid support	License-based, expensive but feature-rich

Final Thoughts



PostgreSQL: Innovation, Freedom, Extensibility 



Oracle: Stability, Enterprise Support, Advanced Features 



Choose based on your needs, not hype!

PostgreSQL 16 Installation

Surendra Panpaliya

Download PostgreSQL Installer

- Visit: <https://www.enterprisedb.com/downloads/postgres-postgresql-downloads>
- Select: **Windows → PostgreSQL 16**

Run the Installer

Double-click
the .exe file

Follow the
Setup Wizard

Select
components:

PostgreSQL
Server 

pgAdmin 4


Stack Builder
(optional)

Set Password & Configure Port

Enter	Enter a superuser password
Configure	Configure Port
Default	Default port: 5432
Change	You can change if needed.

Choose Data Directory



Default: C:\Program
Files\PostgreSQL\16\data



Finish Installation

macOS Installation

- **Use EnterpriseDB Installer**
- Download from:
<https://www.enterprisedb.com/downloads/postgres-postgresql-downloads>
- Run the .dmg installer

Connect to PostgreSQL

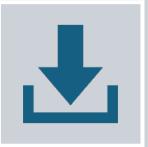
`psql postgres`

Default User:

Username: your system user

No password by default (set manually if needed)

Install pgAdmin 4



Download pgAdmin



Visit: <https://www.pgadmin.org/download/>



Choose Windows / macOS

Run Installer



Install pgAdmin4



Set email and master password (for local access)

Add PostgreSQL Server

Open pgAdmin

Right-click →
Register →
Server

Input:

Name: Local
PostgreSQL

Host: localhost

Port: 5432

Username:
postgres

Password: (your
password from
install step)

Install DBeaver (Universal DB Tool)



Download DBeaver



Visit: <https://dbeaver.io/download/>



Step 2:



Install DBeaver



For Windows: Run .exe installer



For macOS: Run .dmg file, drag to Applications

Connect to PostgreSQL

Open DBeaver

Click: Database
→ New
Connection

Select
PostgreSQL

Fill in:

Host: localhost

Port: 5432

Connect to PostgreSQL



Database: postgres



Username: postgres



Password: your password



Test Connection & Finish

Summary

Tool	Purpose
PostgreSQL 16	Database Server
pgAdmin 4	Web-based GUI for PostgreSQL
DBeaver	Universal SQL Client & ER Diagram Tool

Useful Links



PostgreSQL Docs: <https://www.postgresql.org/docs/16/>



pgAdmin Docs: <https://www.pgadmin.org/docs/>



DBeaver Docs: <https://dbeaver.io/docs/>

Overview of pg_catalog, schemas, roles, and databases

Surendra Panpaliya

Databases in PostgreSQL



Self-contained unit of data.



Each database has its own



schemas, tables,



users (roles), and configurations

Key Points

PostgreSQL is a multi-database system.

Each database is isolated from others

(no cross-database queries directly).

Default databases



postgres → Default working database



template1 → Template for new databases



template0 → Clean template (minimal)

Schemas



Namespace inside a database.



Think of it as a folder to organize



tables, views, functions, etc.

Key Points



Each database can have **multiple schemas**.



Default schema: public

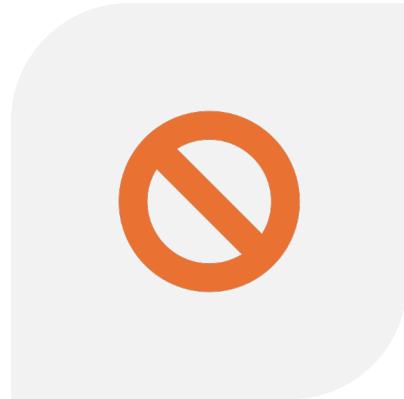


Objects are referenced as:

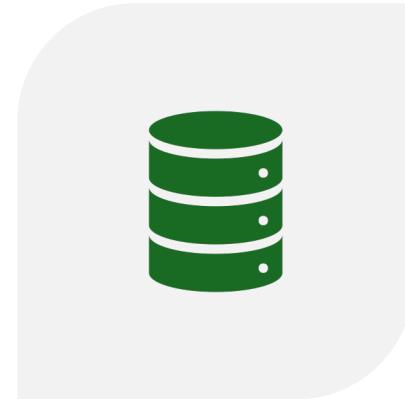


schema_name.object_name

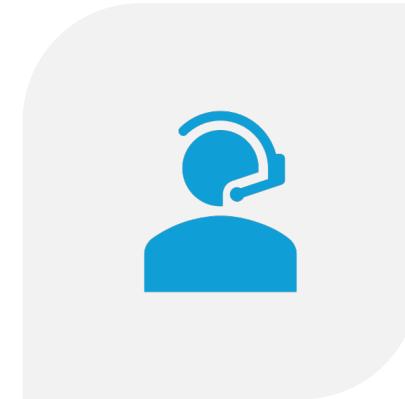
Why Schemas?



AVOID NAME
CONFLICTS



LOGICAL SEPARATION
OF DATA



SUPPORT FOR MULTI-
TENANT APPLICATIONS

Command Examples

-- Create a schema

```
CREATE SCHEMA sales;
```

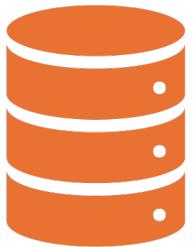
-- Create table inside schema

```
CREATE TABLE sales.orders (id SERIAL PRIMARY KEY, item TEXT);
```

-- Access table

```
SELECT * FROM sales.orders;
```

pg_catalog



A system schema



Stores PostgreSQL

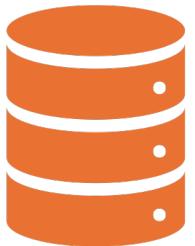


internal metadata.

Contains

Object	Purpose
pg_class	Info about tables, indexes, sequences
pg_user	View of roles with login
pg_roles	Info about all roles
pg_namespace	Info about schemas
pg_tables	List of tables
pg_indexes	List of indexes

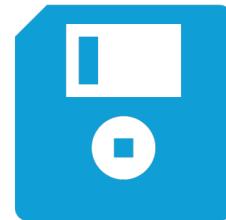
Use Case:



Query database
metadata



Understand internal
PostgreSQL structures



System operations,
backups, monitoring

Example Queries:

-- List all tables using pg_catalog

```
SELECT tablename FROM pg_catalog.pg_tables WHERE  
schemaname = 'public';
```

-- List all roles

```
SELECT rolname FROM pg_catalog.pg_roles;
```

Roles

A Role in PostgreSQL = User Account
+ Group

Roles manage authentication &
authorization.

Types of Roles:

Role	Description
Login Role	Can log in (CREATE ROLE username LOGIN;)
Group Role	Used for permissions, no login

Privileges Controlled By Roles:

CONNECT	CONNECT (to database)
CREATE	CREATE (schemas, tables)
SELECT	SELECT, INSERT, UPDATE, DELETE (on tables)
EXECUTE	EXECUTE (functions)

Role Management Examples



-- Create role with login



```
CREATE ROLE devuser LOGIN PASSWORD 'mypassword';
```



-- Grant role permissions



```
GRANT CREATE ON DATABASE mydb TO devuser;
```

Role Management Examples

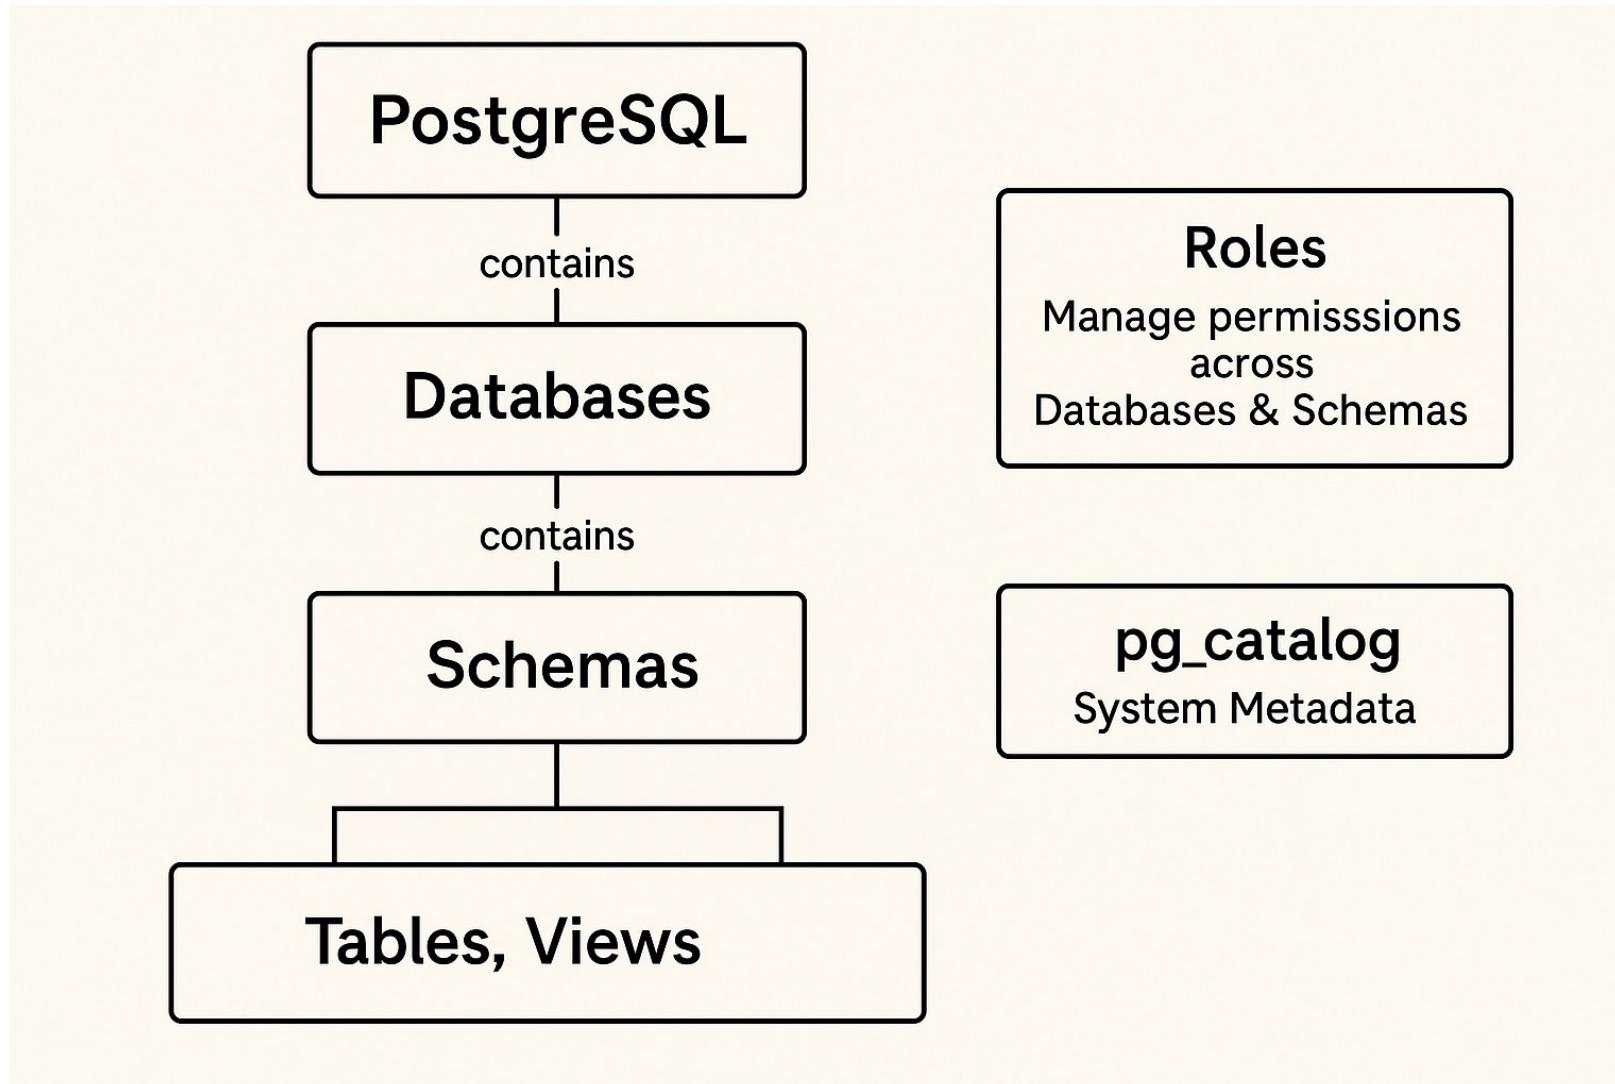
-- Create group role

```
CREATE ROLE analyst;
```

-- Assign user to group

```
GRANT analyst TO devuser;
```

Relationship Diagram



Summary Table

Concept	Purpose
Database	Top-level data container
Schema	Namespace within a database
<code>pg_catalog</code>	Stores metadata about objects
Role	User and permission management

Real-World Analogy

PostgreSQL Concept	Analogy
Database	Apartment Building
Schema	Individual Apartment
Table/View	Furniture inside an apartment
Role	Security Guard controlling access
pg_catalog	Building Management Records



**Thank you for
your support and
patience**

**Surendra Panpaliya
Founder and CEO
GKTCS Innovations**

<https://www.gktcs.com>