



PostgreSQL

**Advanced Postgres**

Surendra Panpaliya

# Schema Design & Data Types

**Surendra Panpaliya**

# Agenda

Tables, schemas, sequences

PostgreSQL vs Oracle data types

SERIAL vs GENERATED ALWAYS AS IDENTITY

# Hands-On

Create customer-product schema

using SERIAL and constraints

# Assignment

Convert an Oracle schema DDL to PostgreSQL DDL

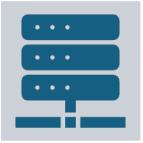
# Tables, schemas, sequences

**Surendra Panpaliya**

# Understanding the Mapping

Oracle Concept	PostgreSQL Equivalent
Schema	Schema
Table	Table
Sequence	Sequence (Auto Number)
Dual Table / FROM DUAL	Use SELECT without FROM or use SELECT 1;

# Example: Create Schemas



Oracle Style:



CREATE USER telecom IDENTIFIED BY password;



Oracle creates a **user-schema** by default.

# Example: Create Schemas



**PostgreSQL Equivalent:**



`CREATE SCHEMA telecom;`



**Note:** In PostgreSQL, `CREATE SCHEMA`



just creates a namespace (not a user).

# PostgreSQL (Preferred way using SERIAL):

```
CREATE TABLE telecom.customers (
    customer_id SERIAL PRIMARY KEY,
    name VARCHAR(100),
    mobile_number VARCHAR(15),
    city VARCHAR(50),
    balance NUMERIC(10,2)
);
```

# Oracle Equivalent:

```
CREATE TABLE telecom.customers (
    customer_id NUMBER GENERATED BY DEFAULT AS IDENTITY
    PRIMARY KEY,
    name VARCHAR2(100),
    mobile_number VARCHAR2(15),
    city VARCHAR2(50),
    balance NUMBER(10,2)
);
```

# Insert Data

```
INSERT INTO telecom.customers (name, mobile_number, city,  
balance)
```

```
VALUES
```

```
('John Doe', '9876543210', 'Pune', 500.00),  
('Jane Smith', '8765432109', 'Mumbai', 1200.00);
```

# Create Plans Table

```
CREATE TABLE telecom.plans (
    plan_id SERIAL PRIMARY KEY,
    plan_name VARCHAR(50),
    monthly_fee NUMERIC(10,2)
);
```

# Create Subscriptions Table (Foreign Keys)

```
CREATE TABLE telecom.subscriptions (
    subscription_id SERIAL PRIMARY KEY,
    customer_id INT REFERENCES telecom.customers(customer_id),
    plan_id INT REFERENCES telecom.plans(plan_id),
    start_date DATE DEFAULT CURRENT_DATE
);
```

# Insert Data into Plans

```
INSERT INTO telecom.plans (plan_name, monthly_fee)  
VALUES ('Basic', 199.00), ('Premium', 499.00);
```

# Query Data with Join

```
SELECT c.name, p.plan_name, s.start_date  
FROM telecom.customers c  
JOIN telecom.subscriptions s ON c.customer_id = s.customer_id  
JOIN telecom.plans p ON s.plan_id = p.plan_id;
```

# View Sequences

List sequences in telecom schema:

```
SELECT sequence_name FROM information_schema.sequences  
WHERE sequence_schema = 'telecom';
```

Check next value of a sequence:

```
SELECT nextval('telecom.customer_seq');
```

# Drop Objects

```
DROP TABLE telecom.subscriptions;  
DROP TABLE telecom.customers;  
DROP TABLE telecom.plans;  
DROP SEQUENCE telecom.customer_seq;  
DROP SCHEMA telecom CASCADE;
```

# Summary

Concept	Oracle	PostgreSQL
Schema	USER/SCHEMA	CREATE SCHEMA
Auto-Increment	IDENTITY	SERIAL or SEQUENCE
Dual Table	FROM DUAL	SELECT 1;
Sequences	CREATE SEQUENCE	Same syntax

# PostgreSQL vs Oracle data types

**Surendra Panpaliya**

# Core Data Type Comparison

Oracle	PostgreSQL	Notes
<b>NUMBER(p,s)</b>	NUMERIC(p,s) / DECIMAL(p,s)	Exact precision, stores as text internally.
<b>NUMBER (no precision)</b>	NUMERIC / DOUBLE PRECISION	NUMERIC for exact; DOUBLE PRECISION for approximate.

# Core Data Type Comparison

Oracle	PostgreSQL	Notes
VARCHAR2(n)	VARCHAR(n)	Functionally identical. PostgreSQL does not distinguish VARCHAR and VARCHAR2.
CHAR(n)	CHAR(n)	Fixed-length, padded with spaces.

# Core Data Type Comparison

Oracle

PostgreSQL

Notes

**CLOB**

**TEXT**

No size limit in PostgreSQL.

**BLOB / RAW**

**BYTEA**

Binary data.

# Core Data Type Comparison

Oracle	PostgreSQL	Notes
DATE	DATE	Both store <b>YYYY-MM-DD</b> . PostgreSQL DATE has no time part.
TIMESTAMP	<b>TIMESTAMP WITHOUT TIME ZONE</b>	Same format, no timezone.

# Core Data Type Comparison

Oracle	PostgreSQL	Notes
<b>TIMESTAMP WITH TIME ZONE</b>	<b>TIMESTAMP WITH TIME ZONE</b>	PostgreSQL handles time zones more strictly.
<b>INTERVAL</b>	<b>INTERVAL</b>	Supports complex date intervals (years, months, days, etc.).

# Core Data Type Comparison

Oracle	PostgreSQL	Notes
<b>BOOLEAN (Oracle 12c+)</b>	<b>BOOLEAN</b>	PostgreSQL has native BOOLEAN. Oracle uses workarounds in older versions.
<b>ROWID</b>	No direct equivalent	Use ctid for internal row address (not for application use).
<b>BFILE</b>	No direct equivalent	Use BYTEA or store file paths.

# Auto-Incrementing IDs

Oracle	PostgreSQL
<b>NUMBER GENERATED BY DEFAULT AS IDENTITY</b>	<b>SERIAL / BIGSERIAL / GENERATED AS IDENTITY</b>

# Special Types

Oracle	PostgreSQL	Use Case
<b>XMLTYPE</b>	XML	XML storage
<b>JSON / JSONB</b>	JSON / JSONB	PostgreSQL natively supports JSON with indexing
<b>USER-DEFINED TYPES (UDT)</b>	COMPOSITE TYPES / DOMAIN	Custom data types

# Function & String Handling

Oracle	PostgreSQL	Notes
SYSDATE	CURRENT_DATE / NOW()	NOW() returns timestamp.
SYSTIMESTAMP	CURRENT_TIMESTAMP	
TO_DATE()	TO_DATE()	Same function name but syntax may vary.
NVL()	COALESCE()	Same functionality.
DECODE()	CASE	PostgreSQL does not have DECODE().

# Data Types Summary Cheat Sheet

Category	Oracle	PostgreSQL
Integer	NUMBER(10)	INTEGER / BIGINT
Float	BINARY_FLOAT / BINARY_DOUBLE	FLOAT / REAL / DOUBLE PRECISION
Character	VARCHAR2 / CHAR	VARCHAR / CHAR
Large Text	CLOB	TEXT
Binary	BLOB / RAW	BYTEA
Date/Time	DATE / TIMESTAMP	DATE / TIMESTAMP
Boolean	Workaround (CHAR(1))	BOOLEAN
Sequence	SEQUENCE	SEQUENCE
JSON	Oracle 21c JSON	JSON / JSONB

# Key Differences to Remember

Oracle	PostgreSQL
<b>VARCHAR2, NVARCHAR2 distinction</b>	Only VARCHAR
<b>NUMBER with arbitrary precision</b>	Use NUMERIC for precision
<b>BLOB/CLOB separate</b>	BYTEA / TEXT handles large data
<b>DATE includes time</b>	DATE is date only in PostgreSQL
<b>JSON less flexible (Oracle pre-21c)</b>	Native JSON support with JSONB indexing

# Customer Billing Table Example

**Oracle:**

```
CREATE TABLE CUSTOMER_BILLING (
    BILL_ID NUMBER GENERATED BY DEFAULT AS IDENTITY PRIMARY
KEY,
    CUSTOMER_ID NUMBER,
    BILL_AMOUNT NUMBER(10,2),
    BILL_DATE DATE
);
```

# Customer Billing Table Example

**PostgreSQL:**

```
CREATE TABLE customer_billing (
    bill_id SERIAL PRIMARY KEY,
    customer_id INTEGER,
    bill_amount NUMERIC(10,2),
    bill_date DATE
);
```

# SERIAL vs GENERATED ALWAYS AS IDENTITY

Surendra Panpaliya

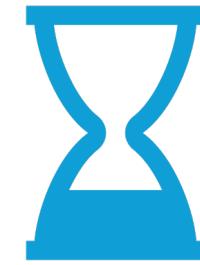
# What is SERIAL?



**SERIAL** is a



**PostgreSQL-specific  
shorthand**



to auto-increment a  
column.

# How it works?

Creates an **INTEGER column**

Automatically creates an **associated SEQUENCE**

Uses `nextval()` to auto-generate values

# Example

```
CREATE TABLE employees (
    emp_id SERIAL PRIMARY KEY,
    name VARCHAR(100)
);
```

# Example

**Behind the scenes:**

PostgreSQL will do this:

```
CREATE SEQUENCE employees_emp_id_seq;
```

# Example

```
CREATE TABLE employees (
    emp_id INTEGER NOT NULL DEFAULT
nextval('employees_emp_id_seq'),
    name VARCHAR(100)
);
```

# What is GENERATED ALWAYS AS IDENTITY?

GENERATED AS  
IDENTITY is

ANSI SQL  
standard syntax

Introduced in  
PostgreSQL 10

# Syntax:

```
CREATE TABLE employees (
    emp_id INT GENERATED ALWAYS AS IDENTITY PRIMARY KEY,
    name VARCHAR(100)
);
```

# Key Differences

Feature	SERIAL	GENERATED AS IDENTITY
Standard Compliance	PostgreSQL-specific	ANSI SQL Standard
Sequence Control	Creates an implicit sequence (manual management needed)	Fully managed internally
Customization	Limited (requires ALTER SEQUENCE)	Supports START WITH, INCREMENT BY, MINVALUE, CYCLE

# Key Differences

Feature	SERIAL	GENERATED AS IDENTITY
DROP TABLE Behavior	Leaves the sequence unless CASCADE is used	Drops the sequence automatically
Insert Override	Can insert custom values by default	<b>ALWAYS:</b> Cannot insert custom values unless OVERRIDING SYSTEM VALUE is used <b>BY DEFAULT:</b> Allows manual inserts

# ALWAYS vs BY DEFAULT

Option	Behavior
<b>GENERATED ALWAYS AS IDENTITY</b>	Auto-generates values and rejects manual insert values (unless overridden)
<b>GENERATED BY DEFAULT AS IDENTITY</b>	Auto-generates values but allows manual inserts if specified

# Example Comparison

SERIAL

```
CREATE TABLE test_serial (
    id SERIAL PRIMARY KEY,
    name TEXT
);
```

Insert custom value:

```
INSERT INTO test_serial (id, name) VALUES (100, 'Custom ID'); -- ✓  
Allowed
```

# IDENTITY

```
CREATE TABLE test_identity (
    id INT GENERATED ALWAYS AS IDENTITY PRIMARY KEY,
    name TEXT
);
```

Insert custom value:

```
INSERT INTO test_identity (id, name) VALUES (100, 'Custom ID');
```

--  ERROR: cannot insert into identity column

# Override behavior

```
INSERT INTO test_identity (id, name)
```

OVERRIDING SYSTEM VALUE

VALUES (100, 'Custom ID'); --  Allowed

# Summary Table

Feature	SERIAL	IDENTITY
SQL Standard	✗	✓
Sequence Name Control	Manual	Automatic
DROP Behavior	Leaves sequence	Drops sequence
Custom Inserts	Allowed	Controlled (ALWAYS/BY DEFAULT)
Preferred For New Systems	🚫 Deprecated in future	✓ Recommended

# Which One Should You Use?

Use Case	Recommendation
New Projects	GENERATED AS IDENTITY (future-proof & standard)
Legacy Support	SERIAL (if maintaining older code)

# Real-World Amdocs Scenario

Oracle	PostgreSQL
<b>NUMBER GENERATED AS IDENTITY</b>	GENERATED ALWAYS AS IDENTITY
<b>SEQUENCE + TRIGGER (older Oracle)</b>	SERIAL or Manual SEQUENCE

# Final Recommendation

Use GENERATED AS IDENTITY

for all new PostgreSQL development

Especially for Oracle to PostgreSQL migrations,

to match modern standards.

# Practice Lab on SERIAL vs IDENTITY in PostgreSQL

Surendra Panpaliya

# Lab Objective

Understand the differences  
between SERIAL and IDENTITY columns  
through practical examples.

# Lab Objective

Concept	Practice
SERIAL usage	✓
GENERATED AS IDENTITY usage	✓
Sequence management	✓
Insert overrides	✓
Drop behavior	✓

# Agenda

## Basic Concepts and SQL Commands

Database, schema, tables, rows, and columns

Data types and constraints

Primary keys and foreign keys

# Agenda

**Basic SQL Commands**

CREATE, INSERT, SELECT, UPDATE, DELETE

Filtering and sorting data

Aggregate function

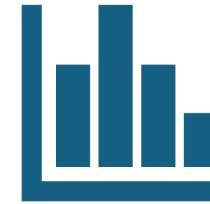
# Database



A PostgreSQL  
database is



Collection of schemas



Store data

# Database

A company database

store data related to employees,

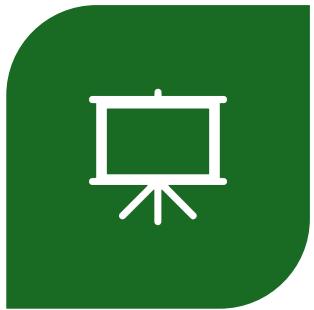
departments, products

`CREATE DATABASE company_db;`

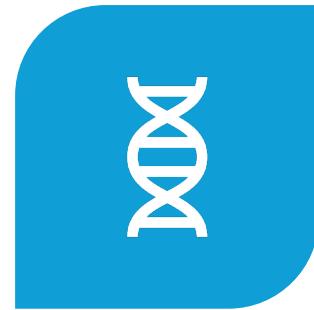
# Schema



LOGICAL CONTAINER  
INSIDE A DATABASE.



HOLDS TABLES,  
VIEWS,



INDEXES,  
SEQUENCES,

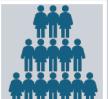


FUNCTIONS

# Example



public (default)



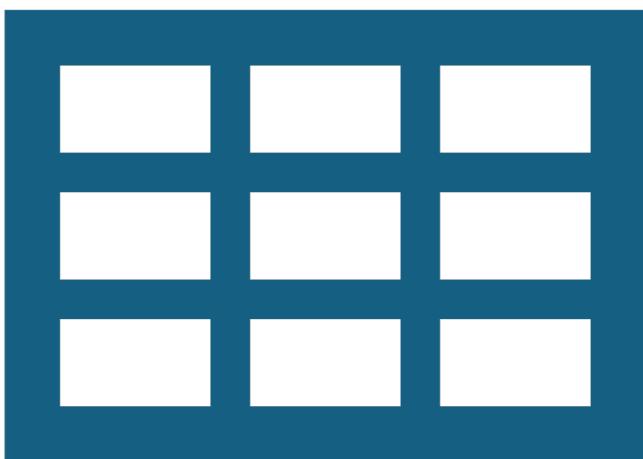
hr (human resources)



sales



CREATE SCHEMA hr;



# Table

Collection of

Related data entries

Consists of

Rows

Columns

# Table Example



An employees table



store details like



employee ID, name,



position, salary

# Creating a Table

```
CREATE TABLE hr.employees (
    employee_id SERIAL PRIMARY KEY,
    first_name VARCHAR(50),
    last_name VARCHAR(50),
    email VARCHAR(100) UNIQUE,
    hire_date DATE,
    salary NUMERIC(8, 2)
);
```

# Creating a Table

`employee_id`:

An integer column that automatically increments with each new row, serving as the primary key.

# Creating a Table

`first_name & last_name:`

A string column with  
a maximum length of  
50 characters.

# Creating a Table

email:

A string column with a maximum length of 100 characters, unique for each row.

# Creating a Table

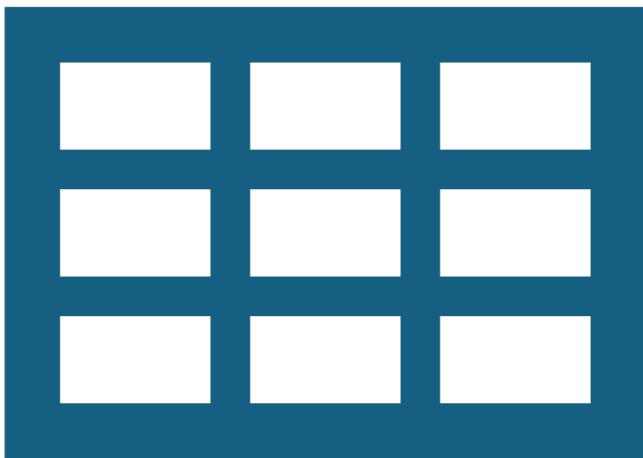
`hire_date:`

A date column.

# Creating a Table

salary:

A numeric column  
allows up to 8 digits,  
including 2 decimal places



# Row

A row (or record)

Represents

Single,

implicitly structured

data item in a table.

# Row Examples

A row in

Employees table

Contain data

about a Single employee.

# Inserting Row into a Table

```
INSERT INTO hr.employees (first_name, last_name,  
email, hire_date, salary)  
VALUES ('Dev', 'Rathi', 'dev.rathi@example.com', '2024-  
07-18', 50000.00);
```

# Column

A column (or field)

a set of data values

of a particular type,

one for each row

of the table

# Column Example

The salary column

in the employee table

stores the salary of

each employee.

# Column

```
SELECT employee_id, first_name, last_name, email,  
hire_date, salary  
FROM hr.employees;
```

# Column

```
ALTER TABLE hr.employees  
ADD COLUMN phone_number VARCHAR(20);
```

# Selecting Data from a Table

```
SELECT employee_id, first_name, last_name, email, hire_date,  
salary FROM hr.employees;
```

# Updating Data in a Table

```
UPDATE hr.employees  
SET salary = 55000.00  
WHERE employee_id = 1;
```

# Deleting Data from a Table

```
DELETE FROM hr.employees  
WHERE employee_id = 1;
```

# **Adding a Column to a Table**

```
ALTER TABLE hr.employees  
ADD COLUMN phone_number VARCHAR(20);
```

# Dropping a Column from a Table

```
ALTER TABLE hr.employees  
DROP COLUMN phone_number;
```

# Dropping a Table

```
DROP TABLE hr.employees;
```

# Dropping a Schema

`DROP SCHEMA hr CASCADE;`

# Dropping a Database

`DROP DATABASE`

# Data Types

Define the kind of data that

Can be stored in each column

Constraints enforce rules on the data.

# Data Types

## SERIAL:

Auto-incrementing integer,

often used for primary keys.

## VARCHAR(n):

Variable-length character string.

# Data Types

**TEXT:**

Variable-length character string

with no specific maximum length.

**INTEGER:** Whole number.

# Data Types

**DATE:**

Calendar date (year, month, day).

**TIMESTAMP:**

Date and time.

# Data Types

**BOOLEAN:**

True or false.

**NUMERIC(p, s):**

Exact numeric of selectable precision.

# Constraints

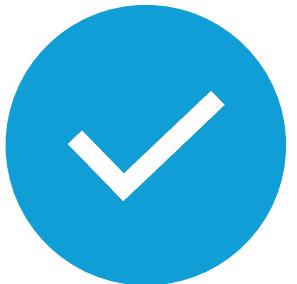
---



Constraints are rules



Applied to columns



To ensure



Data integrity

# Constraints

## **PRIMARY KEY:**

Uniquely identifies each row in a table.

## **FOREIGN KEY:**

Ensures referential integrity between tables.

# Constraints

## NOT NULL:

Ensures that a column cannot have a NULL value.

## UNIQUE:

Ensures that all values in a column are unique.

# Constraints

## CHECK:

Ensures that values in a column  
satisfy a specific condition.

# Numeric Types

## **INTEGER or INT**

Whole numbers.

# Numeric Types

## **SERIAL:**

Auto-incrementing integer

often used for primary keys

# Numeric Types

**NUMERIC(precision, scale):**

Exact numeric values with a specified precision and scale.

**FLOAT, REAL, DOUBLE PRECISION:**

Floating-point numbers.

# **Integer Types**

**smallint:**

Stores 2-byte integer values.

**Range:**

-32768 to 32767.

# **Integer Types**

**integer:**

Stores 4-byte integer values.

**Range:**

-2147483648 to 2147483647.

# **Integer Types**

## **bigint:**

Stores 8-byte integer values.

## **Range:**

-9223372036854775808 to

9223372036854775807

# Integer Types

```
CREATE TABLE numeric_example (
    id serial PRIMARY KEY,
    small_number smallint,
    medium_number integer,
    large_number bigint
);
```

# Integer Types

```
INSERT INTO numeric_example (small_number,  
medium_number, large_number) VALUES (100, 20000,  
3000000000);
```

# **Serial Types**

**serial:**

Auto-incrementing 4-byte integer.

**bigserial:**

Auto-incrementing 8-byte integer.

# Serial Types

```
CREATE TABLE serial_example (
    id serial PRIMARY KEY,
    big_id bigserial
);
```

# **Decimal and Floating-Point Types**

**Decimal or numeric:**

Variable precision number.

**Real:**

4-byte floating point number.

# Decimal and Floating-Point Types

**Double precision:**

8-byte floating point number.

# Decimal and Floating-Point Types

```
CREATE TABLE decimal_example (
    decimal_value decimal(10, 2),
    real_value real,
    double_value double precision
);
```

# Decimal and Floating-Point Types

```
INSERT INTO decimal_example  
(decimal_value, real_value, double_value)  
VALUES (12345.67, 1.23, 123456789.123456);
```

# Character Types

**CHAR(n):**

Fixed-length character string.

# Character Types

**VARCHAR(n):**

Variable-length character string  
with a maximum length of n.

# Character Types

**TEXT:**

Variable-length

character string

with no specific length limit.

# Fixed-Length Character

char(n):

Fixed length, blank-padded

```
CREATE TABLE char_example (
    fixed_length char(10)
);
```

# Fixed-Length Character

```
INSERT INTO char_example (fixed_length)  
VALUES ('abc');
```

# **Variable-Length Character**

**varchar(n):**

Variable length with a limit.

**text:**

Variable unlimited length.

# Variable-Length Character

```
CREATE TABLE varchar_example (
    variable_length varchar(50),
    unlimited_length text
);
```

# Variable-Length Character

```
INSERT INTO varchar_example  
(variable_length, unlimited_length)  
VALUES ('Hello, World!', 'This is a long text field.');
```

# Date/Time Types

## **DATE:**

Calendar date (year, month, day).

## **TIME [ WITHOUT TIME ZONE ]:**

Time of day (hours, minutes, seconds).

# **Date/Time Types**

**TIMESTAMP [ WITHOUT TIME ZONE ]:**

Date and time.

**INTERVAL:**

Time interval.

# **Boolean Type**

**BOOLEAN:**

Logical Boolean

TRUE or FALSE

# Binary Data Types



**BYTEA**



Binary data



byte array

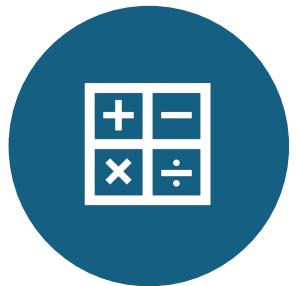
# Array Types



PostgreSQL supports



Arrays of any data type



INTEGER[]



TEXT[]

# Array Types

-- Stores a list of values

```
CREATE TABLE array_example (
    integer_array integer[],
    text_array text[]
);
```

# Array Types

```
INSERT INTO array_example  
(integer_array, text_array)  
VALUES ('{1, 2, 3}', '{"apple", "banana", "cherry"}');
```

# JSON (JavaScript Object Notation)



Stores



JSON data



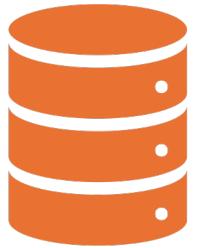
as text.

# JSON Example

```
CREATE TABLE json_example (
    json_data json
);
```

```
INSERT INTO json_example (json_data) VALUES ('{"name": "Dev", "age": 30, "city": "New York"}');
```

# JSONB



Stores JSON data in



Binary format



For faster processing

# JSONB Example

```
CREATE TABLE jsonb_example (
    jsonb_data jsonb
);
```

```
INSERT INTO jsonb_example (jsonb_data) VALUES ('{"name": "Dev", "age": 25, "city": "San Francisco"}');
```

# Constraints

Constraints are rules applied

To table columns

To enforce data integrity

Ensure the accuracy

Reliability of the data

# NOT NULL Constraint

Ensures that a column cannot have a NULL value.

```
CREATE TABLE employees (
    employee_id SERIAL PRIMARY KEY,
    first_name VARCHAR(50) NOT NULL,
    last_name VARCHAR(50) NOT NULL
);
```

# UNIQUE Constraint

Ensures that all values in a column are unique.

```
CREATE TABLE employees (
    employee_id SERIAL PRIMARY KEY,
    email VARCHAR(100) UNIQUE
);
```

# PRIMARY KEY Constraint

A combination of NOT NULL and UNIQUE.  
Uniquely identifies each row in a table.

```
CREATE TABLE employees (
    employee_id SERIAL PRIMARY KEY,
    first_name VARCHAR(50) NOT NULL,
    last_name VARCHAR(50) NOT NULL
);
```

# FOREIGN KEY Constraint

Ensures the referential integrity of the data in one table to match values in another table.

```
CREATE TABLE departments (
    department_id SERIAL PRIMARY KEY,
    department_name VARCHAR(50) NOT NULL
);
```

# FOREIGN KEY Constraint

```
CREATE TABLE employees (
    employee_id SERIAL PRIMARY KEY,
    first_name VARCHAR(50) NOT NULL,
    last_name VARCHAR(50) NOT NULL,
    department_id INTEGER REFERENCES
        departments(department_id)
);
```

# CHECK Constraint

Ensures that all values in a column satisfy a specific condition.

```
CREATE TABLE employees (
    employee_id SERIAL PRIMARY KEY,
    salary NUMERIC(8, 2) CHECK (salary > 0)
);
```

# DEFAULT Constraint

Sets a default value for a column when no value is specified.

```
CREATE TABLE employees (
    employee_id SERIAL PRIMARY KEY,
    hire_date DATE DEFAULT CURRENT_DATE
);
```

# Primary Key

A Column

Combination of columns

Uniquely identifies

each row in a table

# Primary Key

Each table can have

only one primary key,

which can consist of

single or multiple columns.

# Single Column Primary Key

```
CREATE TABLE employees (
    employee_id SERIAL PRIMARY KEY,
    first_name VARCHAR(50),
    last_name VARCHAR(50)
);
```

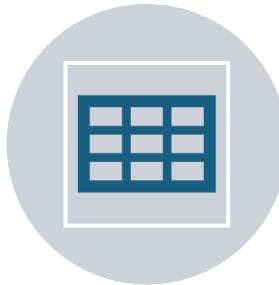
# Composite Primary Key

```
CREATE TABLE project_assignments (
    employee_id INTEGER,
    project_id INTEGER,
    PRIMARY KEY (employee_id, project_id)
);
```

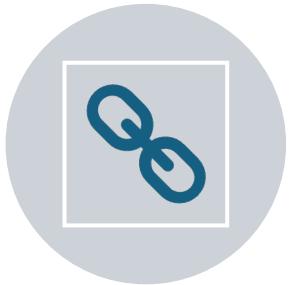
# Foreign Key



A foreign key is



a column (or a  
combination of  
columns)



Establishes a link  
between



data in two tables.

# Foreign Key

Ensures that the value in

Foreign key column(s)

Must match values in

Referenced primary key column(s)

of another table.

# Foreign Key Example

```
CREATE TABLE departments (
    department_id SERIAL PRIMARY KEY,
    department_name VARCHAR(50)
);
```

# Foreign Key Example

```
CREATE TABLE employees (
    employee_id SERIAL PRIMARY KEY,
    first_name VARCHAR(50),
    last_name VARCHAR(50),
    department_id INTEGER,
    FOREIGN KEY (department_id) REFERENCES
departments(department_id)
);
```

# **Foreign Key with ON DELETE CASCADE**

Ensures When a row in

Referenced table is deleted,

All related rows in

Referencing table are deleted.

# Foreign Key with ON DELETE CASCADE

```
CREATE TABLE employees (
    employee_id SERIAL PRIMARY KEY,
    first_name VARCHAR(50),
    last_name VARCHAR(50),
    department_id INTEGER,
    FOREIGN KEY (department_id) REFERENCES
departments(department_id) ON DELETE CASCADE
);
```

# Filtering and Sorting Data

## WHERE Clause

```
SELECT * FROM employees WHERE department_id = 1;
```

- Retrieves all columns and rows
- from the employees table
- where the department\_id is 1.

# Multiple Conditions with AND

```
SELECT * FROM employees  
WHERE department_id = 1 AND salary > 50000;
```

- Retrieves rows from the employees table
- where the department\_id is 1 and
- the salary is greater than 50000.

# Using OR in WHERE Clause

```
SELECT * FROM employees WHERE department_id = 1 OR  
department_id = 2;
```

- Retrieves rows from the employees table
- where the department\_id is either 1 or 2.

# Using IN Operator

```
SELECT * FROM employees WHERE department_id IN (1, 2);
```

- Retrieves rows from the employees table
- where the department\_id is 1 or 2.

# Using LIKE Operator for Pattern Matching

```
SELECT * FROM employees WHERE first_name LIKE 'J%';
```

- Retrieves rows from the employees table
- where the first\_name starts with 'J'.

# Ordering Results with ORDER BY

```
SELECT * FROM employees ORDER BY last_name ASC;
```

- Retrieves all columns and rows
- from the employees table and
- sorts the results by last\_name
- in ascending order (ASC)

# Ordering Results with DESC

```
SELECT * FROM employees ORDER BY salary DESC;
```

- Retrieves all columns and rows
- from the employees table and
- sorts the results by salary in
- descending order (DESC)

# Aggregate Functions

Perform calculations on

Multiple rows

To return

Single value

# Aggregate Functions

## COUNT:

```
SELECT COUNT(*) FROM employees;
```

- Returns the number of rows
- in the employees table

# SUM

```
SELECT SUM(salary) FROM employees;
```

- Returns the sum of
- all salary values
- in the employees table.

# AVG

```
SELECT AVG(salary) FROM employees;
```

- Returns the average
- salary value in
- the employees table.

# MIN

- SELECT MIN(salary) FROM employees;
- Returns the minimum salary value in the employees table.

# MAX

- SELECT MAX(salary) FROM employees;
- Returns the maximum
- salary value in
- the employees table

# Summary



Filtering and sorting data



Essential capabilities in SQL



For retrieving specific subsets of data



Arranging them in a desired order

# Summary



Aggregate functions provide



Powerful tools for calculating



Summary statistics or



Performing calculations



Across multiple rows of data.

# Summary



UNDERSTANDING THESE  
CONCEPTS



CRUCIAL FOR EFFECTIVE  
DATA MANIPULATION



ANALYSIS IN DATABASE  
OPERATIONS

**Surendra Panpaliya**  
**Founder and CEO**  
**GKTCS Innovations**  
<https://www.gktcs.com>

