



PostgreSQL

**Window Functions & Analytics**

Surendra Panpaliya

# Agenda

RANK, DENSE\_RANK, ROW\_NUMBER

LAG, LEAD, FIRST\_VALUE, SUM() OVER

Comparison: Oracle analytic functions vs PostgreSQL

# Hands-On:

Create product sales ranking report with window functions

# Assignment:

Generate per-customer running total using OVER()

# Window functions

---

Provide a way

to perform calculations

across a set of table rows

that are related

to the current row.

# Window functions

---



Useful for  
performing



Aggregations



Calculations



Over subsets  
of data

# What Are Window Functions?

Perform calculations

across rows that are related

to the current row,

without collapsing results

into groups (unlike GROUP BY).

# RANK, DENSE\_RANK, and ROW\_NUMBER

Function	Purpose
<b>ROW_NUMBER()</b>	Assigns a unique number to each row in the partition
<b>RANK()</b>	Assigns the same rank to tied values, with gaps after the tie
<b>DENSE_RANK()</b>	Assigns the same rank to tied values, but without gaps



# WHY are these important in the Telecom Domain?

Telecom datasets are typically large and structured around:

**Customer behavior** (usage, billing, recharges)

**Regional operations** (circles, zones)

**Temporal sequences** (monthly usage, login, recharge)

# **WHY are these important in the Telecom Domain?**

Ranking functions help to:

Identify **top-N customers**

Determine **loyalty tiers**

# WHY are these important in the Telecom Domain?

Track	Track first-time actions
Handle	Handle duplicates
Analyze	Analyze subscription/order history

# WHERE can they be used?

Use Case	Preferred Function
Top 3 data users per telecom circle	RANK or DENSE_RANK
Loyalty program (Gold, Silver, Bronze)	DENSE_RANK
First recharge of every customer	ROW_NUMBER
Track sequence of plan changes	ROW_NUMBER
Deduplicate transactions	ROW_NUMBER
Rank least active users (churn analysis)	RANK

# Comparison Table for Developers

Use Case	Function	Why?
Top N users per circle	RANK()	Allows ties; skips ranks
First transaction/event per group	ROW_NUMBER()	One row per group
Loyalty tiering	DENSE_RANK()	No gaps in ranking
Sequential plan upgrades	ROW_NUMBER()	Ordered tracking
Removing duplicates	ROW_NUMBER()	Keep first; drop rest
Least active customer per circle	RANK()	Useful for churn analysis

# Sample Data Setup

```
CREATE TABLE sales (  
    sale_id SERIAL PRIMARY KEY,  
    customer_name VARCHAR(50),  
    sale_amount NUMERIC(10,2),  
    city VARCHAR(50)  
);
```

# Sample Data Setup

```
INSERT INTO sales (customer_name, sale_amount, city) VALUES  
('Dev', 500, 'Pune'),  
('Dev', 300, 'Pune'),  
('Harish', 400, 'Mumbai'),  
('Harish', 200, 'Mumbai'),  
('Satish', 700, 'Delhi'),  
('Satish', 600, 'Delhi');
```

# RANK Functions

## a) ROW\_NUMBER()

Gives a unique row number within a partition.

```
SELECT customer_name, sale_amount, city,  
ROW_NUMBER() OVER (PARTITION BY city ORDER BY sale_amount  
DESC) AS rn  
FROM sales;
```



# RANK Functions

## b) RANK()

Gives rank with gaps (similar to Oracle RANK()).

```
SELECT customer_name, sale_amount, city,  
       RANK() OVER (PARTITION BY city ORDER BY sale_amount DESC)  
AS rnk  
FROM sales;
```

# RANK Functions

c) DENSE\_RANK()

Gives rank without gaps (like Oracle DENSE\_RANK()).

```
SELECT customer_name, sale_amount, city,  
       DENSE_RANK() OVER (PARTITION BY city ORDER BY  
sale_amount DESC) AS drnk  
FROM sales;
```

# RANK vs DENSE\_RANK vs ROW\_NUMBER

Function	Handles Ties (Same Values)?	Skips Ranks After Ties?	Returns Unique Row Number?	Example Output
RANK()	✓ Yes	✓ Yes	✗ No	1, 2, 2, 4
DENSE_RANK()	✓ Yes	✗ No	✗ No	1, 2, 2, 3
ROW_NUMBER()	✗ No (treats all rows uniquely)	✗ N/A	✓ Yes	1, 2, 3, 4

# Why Are These Important?

Use Case	Best Function	Why?
Ranking top performers (with ties)	RANK() or DENSE_RANK()	Keeps fairness for same scores
Deduplicating records	ROW_NUMBER()	Helps in deleting duplicates (keep row number = 1)
Paginating results	ROW_NUMBER()	Easy offset for pages (e.g., 1–10, 11–20, etc.)
Leaderboards	DENSE_RANK()	No gaps in ranks (used in competitions, contests)
Change detection in partitions	All 3	Track order within partitions

# Dataset (sales2)

---

customer_name	sale_amount	city
Dev	500	Pune
Dev	300	Pune
Harish	400	Mumbai
Harish	200	Mumbai
Satish	700	Delhi
Satish	600	Delhi
Satish	700	Delhi
Satish	600	Delhi

# Use the SQL Query

```
SELECT
    customer_name,
    sale_amount,
    RANK() OVER (PARTITION BY customer_name ORDER BY
sale_amount DESC) AS rank,
    DENSE_RANK() OVER (PARTITION BY customer_name ORDER BY
sale_amount DESC) AS dense_rank,
    ROW_NUMBER() OVER (PARTITION BY customer_name ORDER BY
sale_amount DESC) AS row_num
FROM sales2;
```

# Output of All Three Ranking Functions

---

customer_name	sale_amount	rank	dense_rank	row_num
Dev	500	1	1	1
Dev	300	2	2	2
Harish	400	1	1	1
Harish	200	2	2	2
Satish	700	1	1	1
Satish	700	1	1	2
Satish	600	3	2	3
Satish	600	3	2	4

# **RANK()**

- Ties share the same rank.
- Next rank is skipped.
- For Satish:
- Two 700s → **rank 1**,
- Next is **rank 3** (skipping rank 2).



# ***DENSE\_RANK()***

- Ties share the same rank.
- No rank is skipped.
- For Satish:
- Two 700s → **dense\_rank 1**,
- Two 600s → **dense\_rank 2**

# ***ROW\_NUMBER()***

- Assigns **unique row number**, no tie handling.
- Purely orders within each partition.
- For Satish:
  - First 700 → 1
  - Second 700 → 2
  - First 600 → 3
  - Second 600 → 4

# When to Use What?

---

Use Case	Recommended Function
Eliminate duplicates or pick top-N rows	ROW_NUMBER()
Assign rankings with gaps for ties	RANK()
Assign rankings with no gaps (dense)	DENSE_RANK()
Show position in contests/leaderboards	RANK() or DENSE_RANK()

# LAG() and LEAD()

Window functions

Used to access data from another row

in the result set

without using self-joins.

# WHAT are These Functions?

Function	Description
<b>LAG(column)</b>	Returns the previous row's value
<b>LEAD(column)</b>	Returns the next row's value
<b>FIRST_VALUE()</b>	Returns the first row's value in the window
<b>SUM() OVER()</b>	Computes a running or partitioned total (cumulative or grouped total)

# LAG() and LEAD()

---

Function	Purpose
LAG()	Fetches data from a previous row
LEAD()	Fetches data from a next row

# Syntax



LAG(column\_name, offset, default) OVER (PARTITION BY ...  
ORDER BY ...)



LEAD(column\_name, offset, default) OVER (PARTITION BY ...  
ORDER BY ...)

# Syntax



**column\_name:** the column whose value you want to fetch.



**offset:** how many rows behind (LAG) or ahead (LEAD) to look.  
Default is 1.



**default:** value to return if the target row doesn't exist (optional).



# Example Table: sales

sale_id	customer_name	sale_amount
1	Dev	500
2	Dev	300
3	Dev	200
4	Harish	400
5	Harish	200

# Example 1: Using LAG()

```
SELECT
    customer_name,
    sale_amount,
    LAG(sale_amount) OVER (PARTITION BY customer_name ORDER
BY sale_id) AS previous_sale
FROM sales;
```



# Output

customer_name	sale_amount	previous_sale
Dev	500	NULL
Dev	300	500
Dev	200	300
Harish	400	NULL
Harish	200	400

## Example 2: Using LEAD()

```
SELECT
    customer_name,
    sale_amount,
    LEAD(sale_amount) OVER (PARTITION BY customer_name ORDER
BY sale_id) AS next_sale
FROM sales;
```



# Output

customer_name	sale_amount	next_sale
Dev	500	300
Dev	300	200
Dev	200	NULL
Harish	400	200
Harish	200	NULL

# Real-World Use Cases

Use Case	Function Used
Compare current row to previous sales	LAG()
Detect changes in status or value over time	LAG() / LEAD()
Compute differences between rows (delta)	LAG()
Track next appointment or transaction	LEAD()
Detect gaps in time series data	LAG() / LEAD()

# Advanced Example: Sales Change Detection

```
SELECT  
    customer_name,  
    sale_amount,  
    sale_amount - LAG(sale_amount) OVER (PARTITION BY  
customer_name ORDER BY sale_id) AS change  
FROM sales;
```

# **Advanced Example: Sales Change Detection**

Use this to identify trends in sales increase or decrease.



# Best Practices

- Always use ORDER BY in the OVER() clause to control the sequence.
- Use PARTITION BY when analyzing trends **per customer** or **per group**.
- Use COALESCE(..., 0) if you want to avoid NULLs in your output.

# FIRST\_VALUE() and LAST\_VALUE()

Returns the

**first/last value**

of a window frame

for each row.

# FIRST\_VALUE() and LAST\_VALUE()

```
SELECT
    customer_name,
    sale_amount,
    FIRST_VALUE(sale_amount) OVER (PARTITION BY customer_name
    ORDER BY sale_id) AS first_sale,
    LAST_VALUE(sale_amount) OVER (PARTITION BY customer_name
    ORDER BY sale_id ROWS BETWEEN UNBOUNDED PRECEDING
    AND UNBOUNDED FOLLOWING) AS last_sale
FROM sales;
```

# Note

For LAST\_VALUE(),  
you **must define the frame** correctly,  
else it might return the current row's value  
instead of the true last.

# SUM() OVER (...)



Cumulative or running total



**per partition or**



**over all rows.**

# SUM() OVER (...)

```
SELECT
    customer_name,
    sale_amount,
    SUM(sale_amount) OVER (PARTITION BY customer_name ORDER
BY sale_id) AS running_total
FROM sales;
```

This gives a **progressive sum** for each customer.

# AVG() OVER (...)

Returns a **moving average** or  
group average  
**without collapsing rows.**

# AVG() OVER (...)

```
SELECT  
    customer_name,  
    sale_amount,  
    AVG(sale_amount) OVER (PARTITION BY customer_name) AS  
    avg_sale  
FROM sales;
```

Returns the **same average** for each customer group.



# COUNT() OVER (...)

Counts the number of rows **in the partition**.

# COUNT() OVER (...)

```
SELECT
```

```
    customer_name,
```

```
    sale_amount,
```

```
    COUNT(*) OVER (PARTITION BY customer_name) AS sale_count
```

```
FROM sales;
```

# NTILE(n)

Breaks ordered data into  
**n buckets**  
(quantiles/quartiles).

# NTILE(n)

```
SELECT
```

```
    customer_name,
```

```
    sale_amount,
```

```
    NTILE(2) OVER (PARTITION BY customer_name ORDER BY  
sale_amount DESC) AS quartile
```

```
FROM sales;
```

Breaks sales into 2 groups (top/bottom).

# Real-Life Use Cases

---

Use Case	Function
Running total	SUM() OVER
Cumulative average	AVG() OVER
Finding earliest/latest sale per group	FIRST_VALUE() / LAST_VALUE()
Counting rows per group	COUNT() OVER
Percentile distribution (quantiles)	NTILE(n)

# Best Practices

Use	Use PARTITION BY for grouped analysis
Use	Use ORDER BY for ordered calculations
Combine	Combine with LAG() or LEAD() to compare current vs. past/future values
Use	Use ROWS BETWEEN carefully for correct frame in cumulative analytics

# What Are Analytic Functions?

Oracle	PostgreSQL
<b>Called Analytic Functions</b>	Called Window Functions
<b>Use OVER() clause</b>	Use OVER() clause
<b>Used for ranking, aggregation, previous/next rows</b>	Same

# Syntax Comparison

Feature	Oracle SQL	PostgreSQL SQL
Partition	PARTITION BY	PARTITION BY
Ordering	ORDER BY	ORDER BY
Window Frame	ROWS BETWEEN	ROWS BETWEEN
Aggregate over window	SUM() OVER()	SUM() OVER()



# Syntax Comparison

Feature	Oracle SQL	PostgreSQL SQL
Rank functions	RANK(), DENSE_RANK(), ROW_NUMBER()	Same
Lag/Lead	LAG(), LEAD()	Same
First/Last	FIRST_VALUE(), LAST_VALUE()	Same
Nth Value	NTH_VALUE()	Same

# Function Mapping

Purpose	Oracle	PostgreSQL
Row Number	ROW_NUMBER()	ROW_NUMBER()
Rank	RANK()	RANK()
Dense Rank	DENSE_RANK()	DENSE_RANK()
Running Total	SUM(col) OVER()	SUM(col) OVER()
LAG/LEAD	LAG(), LEAD()	LAG(), LEAD()
First/Last Value	FIRST_VALUE(), LAST_VALUE()	FIRST_VALUE(), LAST_VALUE()
Nth Value	NTH_VALUE()	NTH_VALUE()

# Syntax Example: RANK

Oracle

```
SELECT name, salary,
```

```
    RANK() OVER (PARTITION BY department ORDER BY salary  
DESC) AS rnk
```

```
FROM employees;
```

# Syntax Example: RANK

PostgreSQL

```
SELECT name, salary,
```

```
    RANK() OVER (PARTITION BY department ORDER BY salary  
DESC) AS rnk
```

```
FROM employees;
```

# Window Frame Defaults

Feature	Oracle Default	PostgreSQL Default
<b>ROWS BETWEEN</b>	UNBOUNDED PRECEDING AND CURRENT ROW (For aggregates)	Same
<b>For RANK, ROW_NUMBER</b>	No frame needed	No frame needed

# Partitioning & Ordering

Oracle Example	PostgreSQL Example
<b>PARTITION BY department</b>	PARTITION BY department
<b>ORDER BY salary DESC</b>	ORDER BY salary DESC

# Key Differences

Feature	Oracle	PostgreSQL
Functionality	Almost identical	Almost identical
Syntax	Same	Same
Performance	Highly optimized in Oracle	Efficient in PostgreSQL but tuning may differ
Advanced Windows	MATCH_RECOGNIZE (Oracle 12c+)	Not available in core PostgreSQL

# Use Case Comparison

Use Case	Oracle	PostgreSQL
Top-N per group	RANK(), ROW_NUMBER()	Same
Running total	SUM() OVER()	Same
Time series gaps	LAG(), LEAD()	Same
Partitioned calculations	PARTITION BY	Same



# Summary Table

Analytic Concept	Oracle	PostgreSQL
Rank Functions	✓	✓
Running Totals	✓	✓
LAG/LEAD	✓	✓
Nth Value	✓	✓
Recursion (CONNECT BY)	✓	Use WITH RECURSIVE
Pattern Matching	MATCH_RECOGNIZE	✗ (Use LAG/LEAD + logic)

# Conclusion

Overall	Oracle	PostgreSQL
<b>Analytic SQL Coverage</b>	✓ Advanced	✓ Standard SQL (99% compatible)
<b>Syntax</b>	Nearly identical	Nearly identical
<b>Migration Effort</b>	Minimal for window functions	Minimal

# Hands-On

- Create a Product Sales Ranking Report using Window Functions in PostgreSQL

# Scenario

Task	Window Function
Rank products by sales	RANK()
Handle ties	DENSE_RANK()
Show order of sale entries	ROW_NUMBER()
Calculate total and running total	SUM() OVER()

# Assignment

Generate Per-Customer Running Total  
Using OVER() in PostgreSQL



**Thank you for  
your support and  
patience**

**Surendra Panpaliya**  
**Founder and CEO**  
**GKTCS Innovations**  
<https://www.gktcs.com>